

Internal and Observational Parametricity for Cubical Agda

ANTOINE VAN MUYLDER, KU Leuven, Belgium

ANDREAS NUYTS, KU Leuven, Belgium

DOMINIQUE DEVRIESE, KU Leuven, Belgium

Two approaches exist to incorporate parametricity into proof assistants based on dependent type theory. On the one hand, parametricity translations conveniently compute parametricity statements and their proofs solely based on individual well-typed polymorphic programs. But they do not offer internal parametricity: formal proofs that any polymorphic program of a certain type satisfies its parametricity statement. On the other hand, internally parametric type theories augment plain type theory with additional primitives out of which internal parametricity can be derived. But those type theories lack mature proof assistant implementations and deriving parametricity in them involves low-level intractable proofs. In this paper, we contribute Agda --bridges: the first practical internally parametric proof assistant. We provide the first mechanized proofs of crucial theorems for internal parametricity, like the relativity theorem. We identify a high-level sufficient condition for proving internal parametricity which we call the structure relatedness principle (SRP) by analogy with the structure identity principle (SIP) of HoTT/UF. We state and prove a general parametricity theorem for types that satisfy the SRP. Our parametricity theorem lets us obtain one-liner proofs of standard internal free theorems. We observe that the SRP is harder to prove than the SIP and provide in Agda --bridges a shallowly embedded type theory to compose types that satisfy the SRP. This type theory is an observational type theory of logical relations and our parametricity theorem ought to be one of its inference rules.

CCS Concepts: • **Theory of computation** → **Type theory**.

Additional Key Words and Phrases: cubical type theory, parametricity, structure relatedness principle, Agda

ACM Reference Format:

Antoine Van Muylder, Andreas Nuyts, and Dominique Devriese. 2024. Internal and Observational Parametricity for Cubical Agda. *Proc. ACM Program. Lang.* 8, POPL, Article 8 (January 2024), 32 pages. <https://doi.org/10.1145/3632850>

1 INTRODUCTION

Theorems for free [Wadler 1989] are mathematical statements about polymorphic programs whose validity only depends on a program's type, not its implementation. Such theorems hold in programming languages that prevent polymorphic programs from inspecting their type arguments. This restriction forces polymorphic programs to behave *parametrically*, i.e., apply the same algorithm irrespective of the type they are invoked at.

For example, let us take a purely functional, polymorphic program taking two lists as input and outputting a single list, for an arbitrary type X (we use curly braces to indicate the presence of an implicit argument).

$$p : \forall \{X : \text{Type}\} \rightarrow \text{List } X \rightarrow \text{List } X \rightarrow \text{List } X \quad (1)$$

Authors' addresses: [Antoine Van Muylder](#), KU Leuven, DistriNet, Belgium, antoine.vanmuylder@kuleuven.be; [Andreas Nuyts](#), KU Leuven, DistriNet, Belgium, andreas.nuyts@kuleuven.be; [Dominique Devriese](#), KU Leuven, DistriNet, Belgium, dominique.devriese@kuleuven.be.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

© 2024 Copyright held by the owner/author(s).

ACM 2475-1421/2024/1-ART8

<https://doi.org/10.1145/3632850>

If p is parametric, its range of possible behaviors is considerably limited. Indeed, it cannot branch on concrete values of X like $X = \text{Bool}$ and thus can only use its inputs xs, ys through the `List` interface: the resulting list $p\ xs\ ys$ must be obtained by interleaving, duplicating or omitting values from xs and ys . As a result, such a parametric program p should satisfy the following theorem:

$$\forall (A_0\ A_1 : \text{Type}) (xs\ ys : \text{List } A_0) (f : A_0 \rightarrow A_1) \rightarrow \text{map } f\ (p\ xs\ ys) \equiv p\ (\text{map } f\ xs)\ (\text{map } f\ ys) \quad (2)$$

The theorem holds when p is a list concatenation function, for instance. But in fact, the reasoning above applies for arbitrary parametric implementations of type $\forall\{X : \text{Type}\} \rightarrow \text{List } X \rightarrow \text{List } X \rightarrow \text{List } X$, which all satisfy the theorem. For this reason, the theorem is “free”, i.e. obtained at zero cost.

Several approaches have been developed to study the theoretical aspects of free theorems and enable their use for program verification. The first formal definition of parametricity was given by [Reynolds 1983] for System F, a.k.a. the second-order polymorphic lambda calculus [Girard 1986, 1972; Reynolds 1974]. Reynolds defined parametricity for System F programs as a form of *preservation of relations*. More precisely, he proceeded in two steps. First, he defined a logical relation for System F types, that is, an inductively defined translation mapping any System F type τ into a set-theoretic relation $\llbracket \tau \rrbracket$ between inhabitants of the type. Second, he proved his *abstraction theorem* stating that every inhabitant $x : \tau$ is related to itself $x \llbracket \tau \rrbracket x$. For example¹, the abstraction theorem for $p : \forall\{X : \text{Type}\} \rightarrow \text{List } X \rightarrow \text{List } X \rightarrow \text{List } X$ implies:

$$\begin{aligned} \forall A_0\ A_1\ (R : \text{Rel}(A_0, A_1)) \rightarrow \forall xs_0\ xs_1\ \text{such that } \text{ListRel}_R\ xs_0\ xs_1 \rightarrow \\ \forall ys_0\ ys_1\ \text{such that } \text{ListRel}_R\ ys_0\ ys_1 \rightarrow \text{ListRel}_R\ (p\ xs_0\ ys_0)\ (p\ xs_1\ ys_1) \end{aligned} \quad (3)$$

The relation ListRel_R relates two lists iff they have the same length and their i -th elements are related by R for every i . Observe that Reynolds’s relational definition of parametricity lets us successfully recover the free theorem (2) by setting $R(x, y)$ iff $f\ x \equiv y$ in property (3).

Reynolds’s abstraction theorem is a metatheoretical statement about programs of System F. When verifying polymorphic programs in dependently typed proof assistants like Coq [The Coq development team 2022] or Agda [Agda Development Team 2023], such a metatheoretical statement is unsatisfactory. At best, it guarantees that the free theorem of interest can be added as an axiom in the proof assistant without jeopardizing the logical consistency of the system. The question then arises whether we can prove this free theorem *in* the proof assistant. To answer this question, two approaches have been developed in the literature: parametricity translations on one hand and dependent type theories with internal parametricity on the other.

Parametricity translations [Anand and Morrisett 2017; Bernardy et al. 2012; Keller and Lasson 2012] enhance Reynolds’s logical relation $\llbracket - \rrbracket$ by acting both at the type level and at the term level. On the one hand, types of dependent type theory (DTT), say $T = (X : \text{Type}) \rightarrow \text{List } X \rightarrow \text{List } X \rightarrow \text{List } X$, are mapped to Reynolds’s relational parametricity statement for that type, i.e., their logical relation $\llbracket T \rrbracket$. For T this is (basically) statement (3) but formulated inside DTT, contrary to what can be done for System F. On the other hand, terms of DTT are mapped to a *proof* of the logical relation at their type. For a program $p : T$ it is a proof of (3).

From the point of view of a proof assistant user, parametricity translations are convenient: solely based on an implementation $p : T$, they compute the statement of T ’s free theorem as well as its proof for p . Nonetheless, parametricity translations are not as powerful as they could be. Indeed

¹For the sake of this example we imagine that System F features a `List` type.

they fail to provide proofs of free theorems like the following:

$$\begin{aligned} \text{glob-fthm} : & \forall (p : \{X : \text{Type}\} \rightarrow \text{List } X \rightarrow \text{List } X \rightarrow \text{List } X) \\ & \forall (A_0 A_1 : \text{Type}) (xs\ ys : \text{List } A_0) (f : A_0 \rightarrow A_1) \rightarrow \\ & \text{map } f (p\ xs\ ys) \equiv p (\text{map } f\ xs) (\text{map } f\ ys) \end{aligned} \quad (4)$$

We call such a theorem *global* because the quantification on p is an object-level, or internal quantification (so not metatheoretical). Because of this internal quantification, theorems like `glob-fthm` are known to be logically independent from plain DTT [Booij et al. 2016; Boulier et al. 2017]. Another example of (an indirect consequence of) a global free theorem is the correctness of the Church encoding for the type of booleans $\text{Bool} \simeq (X : \text{Type}) \rightarrow X \rightarrow X \rightarrow X$. Since parametricity translations do not alter the logical expressivity of plain DTT, they cannot possibly validate global free theorems.

An alternative approach that does validate such global free theorems in DTT is the use of *internally parametric type theories* [Cavallo and Harper 2021; Moulin 2016; Nuyts and Devriese 2018; Nuyts et al. 2017]. In such type theories, free theorems can be proven from first principles, even theorems like (4). Those type theories draw inspiration from parametric denotational models of DTT (see e.g. [Atkey et al. 2014]) and augment plain DTT with so-called *parametricity primitives*: additional types, terms and equations (definitional equalities) from which free theorems can be derived. One such parametricity primitive, which appears in all internally parametric DTTs we know of, is the bridge type former (some systems use a different name). It axiomatizes the logical relation $\llbracket T \rrbracket$ at each type T . Internal parametricity then follows from the fact that programs preserve such bridges, similar to how parametric programs preserve relations.

Nonetheless, the cost of the extra logical expressivity granted by internally parametric DTTs is twofold. For one thing, only experimental, unpractical or incomplete implementations of such type theories exist (see e.g. [Cavallo 2020; Nuyts et al. 2017]). More fundamentally, even proving simple free theorems in those systems can be cumbersome. For example, Cavallo and Harper [2021] need about 25 lines of on-paper proof relying on their low-level parametricity primitives to establish that $\text{Bool} \simeq (X : \text{Type}) \rightarrow X \rightarrow X \rightarrow X$.

In this work we contribute a practical proof assistant and associated library combining the two approaches to parametricity in dependent type theory. Concretely, we first contribute the Agda --bridges proof assistant: an extension of the --cubical mode of the Agda proof assistant [Vezzosi et al. 2021]. It implements the parametricity primitives of Cavallo and Harper [2021] (CH), which we have adapted to the cubical type theory underlying Agda --cubical [Cohen et al. 2017]. In order to soundly emulate the substructurality of the bridge type former of CH and their variable-capturing equational theory, we build on the implementation of ticks [Mannaa and Møgelberg 2018] in Agda --guarded [Veltri and Vezzosi 2020, 2023] and let Agda --bridges raise freshness constraints on free variables at typechecking time. Furthermore, Agda --bridges reimplements the equational theory of the Kan operations `hcomp`, `transp` of Agda --cubical with respect to a novel, extended cofibration logic (a.k.a. face logic, see e.g. [Rose et al. 2022] for some background and references). Agda --bridges successfully compiles the full Agda --cubical standard library [Agda Community [n. d.]], offering evidence of its practicality and backwards compatibility. In particular, Agda --bridges conveniently enjoys strong extensionality principles like the univalence and function extensionality theorems provided by Agda --cubical. Within Agda --bridges, we give the first fully formal proofs of several theorems fundamental to internal parametricity, such as relativity [Cavallo and Harper 2021], the relational counterpart of univalence.

Agda --bridges can be used to prove (global) free theorems by hand using its low-level parametricity primitives, as is customary in internally parametric DTTs. However, experience showed

that such direct low-level proofs suffer from several drawbacks. First, familiarity with the CH parametricity primitives is required of the user. Second, these proofs lack compositionality: for example, it is unclear how to reuse a proof like (4) in order to shortcut a parametricity proof at a type built using $T = (X : \text{Type}) \rightarrow \text{List } X \rightarrow \text{List } X \rightarrow \text{List } X$ as a subterm. Third, these proofs are typically long and quickly get intractable as the complexity of the polymorphic target type grows.

To remedy this situation we first observe that the information required to carry out such proofs is exactly captured by a high-level metatheoretical principle which we call *the structure relatedness principle* (SRP). The SRP is a relational, dependent version of the *structure identity principle* (SIP; see e.g. Angiuli et al. [2021b] and Section 6.3). In essence, the SRP simply asserts that for each (dependent) type T , the bridge type of T is equivalent to the logical relation type at T . For a type T like (1), such an equivalence looks like the following, where p_0 and p_1 are polymorphic programs of type T and the left-hand side looks like (3):

$$\begin{aligned} \forall (A_0 A_1 : \text{Type}) (R : A_0 \rightarrow A_1 \rightarrow \text{Type}) &\rightarrow \forall x s_0 x s_1 (x s r : \text{ListRel}_R x s_0 x s_1) \rightarrow \dots \\ &\simeq \text{Bridge}_{(X:\text{Type}) \rightarrow \text{List } X \rightarrow \text{List } X \rightarrow \text{List } X} p_0 p_1 \end{aligned}$$

Note that the SIP asks for similar characterizations, but for path types instead of bridge types.

The SRP at a certain (dependent) type, i.e., a characterization like the above, is enough to derive global free theorems for it. This is the content of our general parametricity theorem `param` (see Section 3.3.2), formulated in Agda --bridges for (dependent) types that satisfy the SRP. It is an internal abstraction theorem asserting that dependent functions preserve or act on logical relations. In practice, proofs of free theorems in Agda --bridges are one-liner invocations of `param`, if the appropriate type has been proven to satisfy the SRP.

Factoring the proof of a free theorem into an SRP proof obligation plus a call to the `param` theorem is already helpful. Such proofs are conceptually easier (with the same computational content) than their low-level counterparts and SRP proofs for simpler types compose to SRP proofs for more complex types. Nonetheless, proving the SRP for dependent types turns out to be challenging. Indeed we observe that some tools that facilitate SIP proofs do not translate to the relational setting: this includes the fundamental theorem of identity types (Theorem 3.5) and the useful fact that the proof-irrelevant fragment of a type can be ignored while proving the SIP for it (Theorem 3.6).

This motivates us to introduce a shallowly embedded domain-specific language (DSL) developed in Agda --bridges, letting the user combine types that validate the SRP. The mere act of writing a type in this DSL, amounts to the construction of an Agda --bridges type that satisfies the SRP, so that `param` can be used straightforwardly. Since it is a tool for composing dependent types, our DSL is in fact a dependent type theory itself, or rather a shallow embedding of a type theory. For the expert reader, our DSL consists of a (raw) category-with-families (CwF) structure on the category of types that satisfy the SRP. Our DSL can be seen as an observational type theory (OTT) [Altenkirch and Kaposi 2015; Altenkirch et al. 2022, 2007; Pujet and Tabareau 2022] of logical relations and our `param` theorem ought to be one of its inference rules. Indeed `param` can be seen as the relational analogue of the *ap* inference rule of OTT stating that open terms act on identifications (equalities). For this reason we call our DSL *relational observational type theory* (ROTT) and say that our `param` theorem is not just internal (to Agda --bridges) but also observational.

ROTT and `param` provide free theorems of practical and theoretical relevance as one-liners in Agda --bridges²: we prove theorem (4); we prove a scheme of Church encodings parametrized by a strictly positive functor; we give the first proof of Reynolds's abstraction theorem for (predicative) System F, using an internally parametric DTT (Agda --bridges) as a metatheory. This was not possible in [Nuyts et al. 2017]. We believe this is the first formal connection between an internally parametric system and Reynolds's relational parametricity. To summarize, our contributions are as follows:

- Agda --bridges: the first practical proof assistant with support for internal parametricity. It implements an adaptation of the internally parametric DTT of Cavallo and Harper [2021] (CH) to the type theory underlying Agda --cubical [Cohen et al. 2017] (CCHM). To faithfully implement the CH theory, Agda --bridges raises freshness constraints at typechecking time; the implementation of this feature builds on the implementation of ticks [Mannaa and Møgelberg 2018] in Agda --guarded [Veltri and Vezzosi 2020, 2023]. Agda --bridges reimplements the `hcomp`, `transp` operations of Agda --cubical with respect to an extended cofibration logic (a.k.a. face logic). It successfully compiles the full Agda --cubical standard library [Agda Community [n. d.]] and thus features the univalence and function extensionality theorems.

- Within Agda --bridges, we provide the first machine-checked proofs of several theorems of fundamental importance for internal parametricity, such as relativity.

- We identify a sufficient condition to obtain free theorems internally, the structure relatedness principle (SRP). We state and prove `param`: a general binary parametricity theorem (abstraction theorem) inside Agda --bridges, formulated for dependent types validating the SRP. Internal free theorems can be obtained out of `param` as one-liners.

- We identify objective reasons why, at a given type, proving the SRP is generally harder than proving the SIP. For these reasons, we provide ROTT: a shallowly embedded type theory for obtaining SRP proofs by merely writing a type, in the spirit of parametricity translations. Our `param` theorem ought to be one of the inference rules of ROTT. Internal parametricity proofs performed using ROTT and its `param` rule are user-friendly, compositional and concise.

- We demonstrate the use and generality of Agda --bridges, ROTT and `param` on practically and theoretically relevant examples.

Outline. This paper is structured as follows. In Section 2 we present the implementation and typing rules of Agda --bridges as well as its core theorems. More precisely: Section 2.1 is a brief reminder about cubical type theory and Agda --cubical; Section 2.2 discusses what makes Agda --bridges and the CH type theory substructural type systems. Sections 2.3, 2.4, 2.5 introduce the `BridgeP`, `extent` and `Gel` parametricity primitives; Section 2.6 and Section 2.7 prove core or basic theorems using Agda --bridges. The discussion about the Kan operations `transp`, `mhcomp` of Agda --bridges and their custom cofibration logic is rather postponed to Section 5 as they are not needed for the free theorems we present. In Section 3 we present a framework developed in Agda --bridges that provides abstractions to write user-friendly, modular and concise proofs of free theorems². More precisely: In Section 3.1 we discuss the structure relatedness principle (SRP); In Section 3.2 we identify obstructions to the SRP; In Section 3.3 we present ROTT and its `param` rule. We use ROTT and apply `param` on various examples in Section 4. We conclude with related work in Section 6.

Conventions and notations. We use the term “logical relation” for (1) a relation R between structured types, compatible with the structure (e.g. a structure-preserving relation $R : M_0 \rightarrow M_1 \rightarrow \text{Type}$ between two monoids); (2) a proof of relatedness for such a structured relation (e.g. a proof $\text{pf} : R\ m_0\ m_1$ for some $m_0 : M_0, m_1 : M_1$); and (3) (more rarely) a relation defined by induction on the formation of types (like $\llbracket - \rrbracket$ above). Our choice to use the same term for (1) and (2) mirrors the situation for bridges. We assume that the reader is somewhat familiar with intensional dependent type theory and proof assistants based on it. The term “free theorem” designates consequences of relational parametricity for potentially non-graph relations R , at any given polymorphic type. For us, the equational theory of a type theory is the collection of its undirected definitional equalities and it includes β and η rules. We typically use Agda syntax to write types, programs and proofs. We ignore writing universe levels. We use curly braces for implicit arguments. Variable binding may

²All of our theorems are part of an Agda --bridges library <https://github.com/antoinevanmuylder/bridgy-lib>.

be denoted with $\lambda x. f$, $\lambda x \rightarrow f$ (Agda syntax) or sometimes just $x.f$. Logical relations between e.g. x_0 and x_1 are typically denoted with a postfix r notation xr . Bridges and paths between x_0 and x_1 are rather denoted with a double-letter notation xx . The ε symbol systematically ranges in $\{0, 1\}$.

2 THE INTERNAL PARAMETRICITY OF AGDA BRIDGES

Agda --bridges is a proof assistant extending the Agda --cubical proof assistant [Vezzosi et al. 2021]. Accordingly, the type theory that Agda --bridges implements, i.e., its primitives and their equational theory, is an extension of the type theory that Agda --cubical implements (called CCHM in the literature [Cohen et al. 2017]). A reminder about Agda --cubical appears in Section 2.1.

The type theory of Agda --bridges is an adaptation of the internally parametric DTT of Cavallo and Harper [2021] (referred to as the CH theory). In other words, Agda --bridges implements the primitives and equations specified by the CH theory (we mostly keep the same names), or rather relatively close variants. The CH theory is not entirely standard as it is a *substructural* (alternatively “affine”) type theory. Indeed, most of its parametricity primitives have typing rules, including operational semantics, that can only be used if certain conditions on free variables are satisfied. This is discussed in Section 2.2. The first main difference between the Agda --bridges and CH theories lies in how they both handle substructurality. Our solution, *freshness typechecking constraints* on free variables, is discussed in Section 2.3. Note that the other main difference w.r.t. the CH theory is that both type theories extend different kinds of cubical type theories. The latter difference is rather discussed in Section 5.

In the CH theory, the bridge type former is postulated to represent *logical relations*: relations between types that respect their structure, or proofs of relatedness under such relations. Concrete examples of logical relations will appear below or can be consulted in [Hermida et al. 2013], for example. To ensure that bridges do uniquely correspond to logical relations, additional primitives called *extent* and *Gel* are postulated by the theory. Internal parametricity then refers to the fact that all programs preserve bridges, which are in one-to-one correspondence with logical relations. Accordingly, Agda --bridges features primitives `BridgeP`, `extent` and `Gel` whose implementation and rules are explained in Sections 2.3, 2.4, 2.5. Occurrences of these primitives in a program or type may generate freshness constraints at typechecking time.

In Section 2.6 and Section 2.7 we use the above primitives to derive within Agda --bridges core theorems for internal parametricity as well as the free theorem $\text{Bool} \simeq (X : \text{Type}) \rightarrow X \rightarrow X \rightarrow X$ in a low-level style.

2.1 The Cubical Fragment of Agda Bridges

Agda --cubical is an implementation of cubical type theory (CCHM) on top of the Agda proof assistant. Overall, cubical type theory modifies standard intensional type theory in several respects. First, it treats proofs of equality as if they were topological *paths* (see Section 2.1.1). Second, it features language primitives that let it realize *univalence* as a theorem (see Section 2.1.3). The latter property characterizes type equality as type equivalence (i.e. having a “bijection”) and constitutes a defining aspect of *homotopy type theory* (HoTT) [Program 2013]. Earlier instances of HoTT assume univalence as an axiom instead. Third, these primitives include the so-called Kan operations called `transport`, `hcomp` in Agda --cubical. Our adaptation of the Kan operations is rather discussed in Section 5. Lastly, as an instance of HoTT, cubical type theory defines higher inductive data types (HITs) which we do not discuss in this work. We now remind the reader about paths, equivalences, univalence and other extensionality principles.

2.1.1 Paths. A major difference between Agda and Agda --cubical lies in how each system treats propositional equality. By contrast with the inductively defined equality types \equiv of standard

intensional type theory, Agda --cubical defines equality in terms of a special postulated type called the *path interval* denoted `I` and equipped with two constants `i0, i1` representing the interval's endpoints. A proof of equality of $a_0, a_1 : A$ is then by definition a function $aa : I \rightarrow A$ such that $aa\ i0 = a_0$ and $aa\ i1 = a_1$ (definitionally). Proofs of equality are commonly called *paths* in the context of cubical type theory. In order to validate basic lemmas about path equality, `I` carries operations `~, ^, v` for manipulating path variables. For instance, the map `~ : I → I` inverts the two endpoints and can be used to prove symmetry of path equality: $\lambda aa\ i \rightarrow aa\ (\sim i) : a_0 \equiv a_1 \rightarrow a_1 \equiv a_0$.

To explain what are paths $p : a_0 \equiv_A a_1$ when A is constituted from dependent types (e.g. A is a Σ -type), an essentially unavoidable notion of *heterogeneous*, or *dependent* path emerges. For this reason Agda --cubical features types `PathP` of dependent paths. Given a line of types, that is, a function $A : I \rightarrow \text{Type}$ and terms $a_0 : A\ i0$ and $a_1 : A\ i1$, one can form the type `PathPA a0 a1 : Type` of heterogeneous or dependent paths between a_0 and a_1 . While non-dependent paths $bb : b_0 \equiv_B b_1$ are functions $I \rightarrow B$ with definitionally fixed endpoints b_0, b_1 , *dependent* paths $aa : \text{PathP}_A\ a_0\ a_1$ are *dependent* functions $(i : I) \rightarrow A\ i$ with definitionally fixed endpoints a_0, a_1 . In fact, the type $b_0 \equiv_B b_1$ of non-dependent paths in B is defined as the type `PathPλi→B b0 b1` of paths over a constant line $(\lambda i \rightarrow B : I \rightarrow \text{Type})$. Note that we will use the notation $i. A\ i := \lambda i \rightarrow A\ i$ to refer to lines of types in inlined code.

2.1.2 Equivalences. In Agda --cubical, an equivalence is a function $f : A_0 \rightarrow A_1$ satisfying the `isEquiv : (A0 → A1) → Type` predicate. The type $A_0 \simeq A_1$ of equivalences between A_0 and A_1 is defined as $\Sigma[f \in (A_0 \rightarrow A_1)] \text{isEquiv } f$. The exact definition of `isEquiv` can be consulted in [Program 2013] e.g., and we rather indicate here a sufficient condition for `isEquiv f` (our equivalences are always built this way). In order to prove `isEquiv f` it is sufficient to build an inverse for $f : A_0 \rightarrow A_1$, that is, a function $g : A_1 \rightarrow A_0$ such that $\forall a_0 \rightarrow g(f\ a_0) \equiv a_0$ and $\forall a_1 \rightarrow f(g\ a_1) \equiv a_1$. Note also that in order to prove that two equivalences $e_0, e_1 : A_0 \simeq A_1$ are equal, it suffices to prove that their underlying functions are equal.

2.1.3 Univalence and Other Extensionality Principles. It is commonplace in proof assistants to be faced with situations where one needs a more concrete representation of an equality type $a_0 \equiv_A a_1$ when the specific shape of A is known. For example, in order to prove that two (perhaps dependent) functions are equal it should be sufficient to prove that they are pointwise equal. In Agda --cubical, this principle admits a simple proof and is realized as an equivalence `funextEquiv : ((a : A) → f0 a ≡B a f1 a) ≃ (f0 ≡(a:A)→B a f1)`.

Such characterizations of equality types $a_0 \equiv_A a_1$ at specific types A are called *extensionality principles*. Interestingly, Agda --cubical is expressive enough to guarantee the validity of similar extensionality principles for all primitive type formers, which we list here. The extensionality principle for Π (or \rightarrow in the non-dependent case) is `funextEquiv`. The principle for Σ (and \times) characterizes the path type $(p_0 \equiv_{\Sigma[a \in A] B a} p_1)$ as $\Sigma[aa \in (p_0.\text{fst} \equiv_A p_1.\text{fst})] \text{PathP}_{i.B(aa\ i)} (p_0.\text{snd}) (p_1.\text{snd})$. Note that `.fst, .snd` extract values from (dependent) pairs. Similarly, extensionality principles for specific data and (even coinductive) record types can always be proven, and there also exists such a principle for the path type itself. Lastly, arguably the most important yet subtle primitive extensionality principle is *univalence*, which asserts that two types are equal if and only if they are equivalent, i.e., `univalence : (A0 ≃ A1) ≃ (A0 ≡Type A1)`. As all the other extensionality principles, univalence is obtained as a theorem in Agda --cubical.

2.2 Substructurality

As explained above, the CH theory is dubbed substructural (or alternatively “affine”) because most of its parametricity primitives have typing rules, including operational semantics, that can only be used if certain conditions on free variables are satisfied. Reasons why these restrictions exist in

the first place appear later in this subsection. The CH theory manages to express such restrictions on free variables by using a *context restriction* operation $\Gamma \mapsto \Gamma \setminus x$ acting on contexts. Typically, restricted contexts $\Gamma \setminus x$ are strictly smaller than Γ and appear in the premises of certain typing rules of CH. Note that having premises featuring smaller contexts is not standard in dependent type theory. We now exemplify context restriction by looking at how the bridge type former is defined in the CH theory.

Substructurality in CH. Recall that the CH theory is an extension of cubical type theory. The first type former presented in CH is the bridge type former. The definition of the bridge type former is similar to that of the path type former as it is also based on the presence of an abstract interval BI called the *bridge interval*. The type BI is assumed to be equipped with two endpoints $\text{bi0}, \text{bi1} : \text{BI}$ but no other operations (like \sim, \wedge, \vee in the case of I, see Section 2.1). This implies that a term $\Gamma \vdash x : \text{BI}$ in an arbitrary context is in fact always either a bridge variable $(x : \text{BI}) \in \Gamma$ or an endpoint $\text{bi0}, \text{bi1}$.

Similar to the path case, the bridge introduction rule (at a closed type A , say) lets us build a bridge $\Gamma \vdash \lambda x. aa : \text{Bridge}_A a_0 a_1$ if a term $\Gamma, x : \text{BI} \vdash aa : A$ is provided. However, the bridge and path type formers feature different elimination rules which we reproduce here:

$$\frac{\Gamma \vdash i : \text{I} \quad \Gamma \vdash aa : \text{Path}_A a_0 a_1}{\Gamma \vdash aa i : A} \quad \frac{\Gamma \vdash x : \text{BI} \quad \Gamma \setminus x \vdash aa : \text{Bridge}_A a_0 a_1}{\Gamma \vdash aa x : A} \quad \text{BDG-ELIM-CH}$$

As can be seen on the right, a bridge aa can be applied to a bridge interval term $\Gamma \vdash x : \text{BI}$ only if aa typechecks in a different, *restricted* context $\Gamma \setminus x$. This restriction operation $\Gamma \mapsto \Gamma \setminus x$ is defined by structural induction on the context. Intuitively, if x is a variable (so not $\text{bi0}, \text{bi1}$) the context $\Gamma \setminus x$ is obtained from Γ by deleting the $x : \text{BI}$ variable from Γ , as well as all those variables “strictly to the right” of x in Γ (whose type could legally contain x). This intuition is made formal in Section 2.3. In particular, the term $\lambda x. \text{sq } x x$ that takes the diagonal of a square of bridges (i.e. a bridge between bridges), is ill-typed. The constraint appearing in BDG-ELIM-CH can be summarized by saying that bridges are only allowed to consume variables $x : \text{BI}$ that are “fresh”. Alternatively one can say that bridges, or the interval BI itself, are substructural, or affine.

The bridge elimination rule is one example of how context restriction is used to make BI substructural. In fact the majority of the rules appearing in the parametricity fragment of the CH theory feature restricted contexts. But knowing that context restriction can be used to enforce substructurality does not explain why the theory is substructural in the first place.

There are essentially two reasons for this, we refer to the CH article for more details. First, CH show that their substructural type theory admits a denotational model (in bicubical sets, see their article). This means that their type theory is logically consistent (relative to Set theory). In other words we can be sure that no logical contradiction can be derived using their primitives, an important requirement for a logic or a proof assistant. Second, several parametricity primitives of the theory (extent, Gel and the Kan operations) present a non standard *variable-capturing* equational theory that is well defined solely thanks to the substructurality of BI. For instance the β -rule of extent operates by identifying a free bridge variable $x : \text{BI}$ in an appropriate argument m and by capturing the variable x in m , yielding an overall term with $\lambda x. m$ as a subterm. More information about capturing appears in Section 2.4.

Substructurality in Agda --bridges. Neither normal dependent Agda functions nor paths of Agda --cubical present an elimination rule with a restricted context. Accordingly, to minimize the implementation work and maximize the reuse of existing Agda infrastructure we do not use a context restriction operation: the premises of our rules do not have smaller contexts compared to their conclusion. For instance, eliminating a bridge aa at a bridge variable $\Gamma \vdash x : \text{BI}$ only requires aa to typecheck in Γ (not $\Gamma \setminus x$). To remain sound w.r.t. the CH theory, Agda --bridges compensates by raising appropriate constraints on free variables when typechecking a rule featuring a premise

where $\Gamma \setminus x$ would appear in the CH theory. An example of such a rule is BDG-ELIM-CH. We call these constraints *freshness typechecking constraints*. Their definitions are given in Definition 2.1, 2.2. Such freshness constraints can be raised on different occasions during typechecking, which we list here.

- We are typechecking a bridge application $\Gamma \vdash aa\ x$.
- We are typechecking an affine function elimination $\Gamma \vdash f\ x$.
- We are computing (so reducing or comparing terms) and need capturing to occur.

Note that affine functions are exactly like bridges, but without fixed endpoints. Affine functions will be used to express the type of `extent`, for example. The first two cases are similar. If the raised constraint is found satisfiable, typechecking can continue and perhaps succeed. Else, a typechecking error occurs. In the third case, a freshness constraint called *semi-freshness* is raised. If it is found satisfiable computing goes on. Else, computing halts (no further reduction, or a failed comparison).

We now present the primitives of Agda --bridges: their typing rules and equational theory, details about their implementation as well as core theorems they guarantee. Following the above, several of these typing rules have premises that are (semi-) freshness constraints. Recall that all the theorems we obtain using Agda --bridges are available in our accompanying library.

2.3 Affine Functions and Bridges

First of all, Agda --bridges postulates the existence of a bridge interval type `BI` equipped with two endpoints `bi0, bi1 : BI` and no further operations. Next, we define the type former of affine functions. Bridges will essentially be affine functions with definitionally fixed endpoints.

2.3.1 Affine Functions. Affine functions are implemented as normal Agda dependent functions but their domain is `BI` and it carries what is called a tick annotation (building on [Veltri and Vezzosi 2020, 2023]). The type of non-dependent affine functions with codomain C is denoted $(@tick\ x : BI) \rightarrow C$ or $@BI \rightarrow C$. We call an affine function $A : (@tick\ x : BI) \rightarrow Type$ an (affine) line of types and such lines are sometimes denoted $x. Ax$.

Given a line $A : (@tick\ x : BI) \rightarrow Type$ one can form the type of dependent affine functions over A denoted $(@tick\ x : BI) \rightarrow Ax$. Compared to normal dependent functions, the tick annotation has the net effect of raising an additional freshness constraint while typechecking the application of a function $f : (@tick\ x : BI) \rightarrow Ax$ to a bridge variable $(x : BI)$ in a given context. That is to say, the following typing rule is implemented.

$$\frac{\Gamma \vdash f : (@tick\ x : BI) \rightarrow Ax \quad (x : BI) \in \Gamma \quad \text{fresh}(f, x)}{\Gamma \vdash f\ x : Ax} \text{ TICK-APP}$$

The other rules of $(@tick\ x : BI) \rightarrow Ax$ are those of normal dependent functions. Notice how this time the context in which f typechecks is not restricted. The constraint on free variables implied by the context restriction operation of (2.2) is instead expressed as a typechecking side condition denoted $\text{fresh}(f, x)$ (understood to live at the same context Γ than f). We call the latter side condition a freshness constraint and it is defined as the following decidable condition on f, x .

Definition 2.1 (Freshness constraint). Setting $\Gamma = \Gamma_1, (x : BI), \Gamma_2$, the freshness constraint $\text{fresh}(f, x)$ is satisfied if for every free variable v of f one of the following holds:

- v appears in Γ_1 (i.e., $v \in \Gamma_1$)
- v appears in Γ_2 and is a path or a bridge variable³, i.e., $(v : I) \in \Gamma_2$ or $(v : BI) \in \Gamma_2$.

We also adopt the convention that `bi0, bi1` are always fresh for any term f .

³Additionally, v can be a variable witnessing the truth of a face constraint (Section 5.2) that does not mention x .

In more mundane terms, f and x satisfy the freshness constraint if f does not mention x as a free variable, nor any term variable (i.e. not a path/bridge variable) declared after x in the context Γ . In what follows the presence of a tick annotation will sometimes be abbreviated to an $@$ symbol, or omitted. To remain sound w.r.t. CH, all functions must use **Bl** with this annotation.

We now explain the **BridgeP**, **extent** and **Gel** type formers. The typing rules implemented for those primitives are provided in Fig. 1. We make a few general remarks about these rules. The ε symbol systematically ranges over the indices 0, 1. An equality judgment $\Gamma \vdash a = b : A$, or just $a = b$ when the context and types are clear, signifies that “the convertibility algorithm of the Agda --bridges typechecker concludes that a and b are definitionally equal”. A reduction judgment $\Gamma \vdash a \mapsto b : A$, or just $a \mapsto b$ signifies that “the reduction algorithm of Agda --bridges finds that a reduces to b ”. More precisely, Agda uses weak head normalization and we denote the result of this algorithm on a by $\text{red } a$. To be more precise about our implementation, some of the rules of Fig. 1 contain occurrences of reduction $\text{red } a$. It is important to know that red is not a type-theoretic primitive but rather a hint that the corresponding argument should be reduced when firing the rule in question. In other words, occurrences of red in the rules must merely be seen as informative decorations providing more details about the implementation. The rules also present occurrences of variable captures $\langle x \rangle a$. Again, capturing is not an object-level program but rather a metatheoretical algorithm applying various substitutions and raising freshness constraints under the hood (see Section 2.4).

2.3.2 Bridges. As hinted above, the rules of **BridgeP** are essentially a replication of the rules of $(@x : \text{Bl}) \rightarrow A x$. According to the formation rule, and given an affine line of types $A : (@x : \text{Bl}) \rightarrow \text{Type}$ as well as two endpoints $a_0 : A \text{ bi0}$ and $a_1 : A \text{ bi1}$ one can form the type **BridgeP** $A a_0 a_1 : \text{Type}$ of heterogeneous or dependent bridges between a_0 and a_1 over the line A . We sometimes write the line A as an index as in **BridgeP** _{A} $a_0 a_1 : \text{Type}$. Given a type $A : \text{Type}$ and two inhabitants a_0, a_1 , the type of non dependent bridges between a_0, a_1 is defined as **Bridge** _{A} $a_0 a_1 = \text{BridgeP}_{x.A} a_0 a_1$. Furthermore, eliminating a bridge at a bridge variable can only be done if a freshness constraint is satisfied, similar to the **TICK-APP** rule of Section 2.3.1. Lastly, when typechecking a bridge $\Gamma \vdash aa : \text{BridgeP}_A a_0 a_1$, the typechecker will verify that $aa \text{ bi}\varepsilon$ and a_ε are definitionally equal. Conversely, the expression $aa \text{ bi}\varepsilon$ will reduce to a_ε for a typechecked bridge.

So far only one example of a bridge can be built in an empty context. If A is a closed type inhabited by a , one can build the reflexivity bridge at A as $\lambda x. a : \text{Bridge}_A a a$. To build examples of bridges different from reflexivity at function types $(a : A) \rightarrow B a$, the **extent** primitive is introduced.

2.4 The Extent Primitive

We now turn our attention to the rules and implementation of the **extent** parametricity primitive. The rules of **extent** are displayed in Fig. 1. The sole purpose of **extent** is to guarantee the validity of a certain *relational* extensionality principle: a principle akin to **funextEquiv** but characterizing bridges between functions rather than paths. We call this principle **extentEquiv** and it reads as:

$$((a_0 a_1 : A)(aa : \text{Bridge}_A a_0 a_1) \rightarrow \text{BridgeP}_{x.B(aa x)}(f_0 a_0)(f_1 a_1)) \simeq \text{Bridge}_{(a:A) \rightarrow B a} f_0 f_1$$

This equivalence asserts that two functions f_0, f_1 are related by a bridge if and only if they are pointwise related, i.e., they map related inputs to related outputs, where “related” signifies the existence of an appropriate bridge. In other words it makes the bridge type former behave as a type former of logical relations, at $(a : A) \rightarrow B a$. Note that there exists a slight generalization of this principle characterizing the type of heterogeneous bridges **BridgeP** _{$x.(a:A x) \rightarrow B x a$} $f_0 f_1$ instead, which we still call **extentEquiv**.

$$\begin{array}{c}
\frac{\Gamma \vdash A : @BI \rightarrow \text{Type} \quad \Gamma \vdash a_\varepsilon : A \text{ bi}\varepsilon}{\Gamma \vdash \text{BridgeP}_A a_0 a_1 : \text{Type}} \text{BDG-FORM} \quad \frac{\Gamma \vdash aa : (@x : BI) \rightarrow Ax \quad \Gamma \vdash aa \text{ bi}\varepsilon = a_\varepsilon}{\Gamma \vdash \lambda x. aa x : \text{BridgeP}_A a_0 a_1} \text{BDG-INTRO} \\
\\
\frac{\Gamma \vdash aa : \text{BridgeP}_A a_0 a_1 \quad \Gamma \vdash x : BI \quad \text{fresh}(aa, x)}{\Gamma \vdash aa x : Ax} \text{BDG-ELIM} \quad \frac{\Gamma \vdash aa : \text{BridgeP}_A a_0 a_1}{\Gamma \vdash aa \text{ bi}\varepsilon \mapsto a_\varepsilon} \text{BDG-}\partial \\
\\
\frac{\Gamma \vdash aa_\varepsilon : \text{BridgeP}_A a_0 a_1 \quad \Gamma, @x : BI \vdash aa_0 x = aa_1 x : Ax}{\Gamma \vdash aa_0 = aa_1 : \text{BridgeP}_A a_0 a_1} \text{BDG-}\eta \quad \frac{\Gamma \vdash aa : (@x : BI) \rightarrow Ax \quad \Gamma \vdash y : BI \quad \text{fresh}(aa, y)}{\Gamma \vdash (\lambda x. aa x) y \mapsto aa y : Ay} \text{BDG-}\beta \\
\\
\frac{\Gamma \vdash A : @BI \rightarrow \text{Type} \quad (x : BI) \in \Gamma \quad \Gamma \vdash a : Ax \quad \text{sfresh}(a, x)}{\Gamma \vdash \langle x \rangle a : (@y : BI) \rightarrow Ay \quad \text{fresh}(\langle x \rangle a, x) \quad \Gamma \vdash (\langle x \rangle a) x = a : Ax} \text{CAP} \\
\\
\frac{\Gamma \vdash A : @BI \rightarrow \text{Type} \quad \Gamma \vdash B : (@x : BI)(a : Ax) \rightarrow \text{Type} \quad \Gamma \vdash f_\varepsilon : (a_\varepsilon : A \text{ bi}\varepsilon) \rightarrow B \text{ bi}\varepsilon a_\varepsilon \quad \Gamma \vdash fr : (a_0 : A \text{ bi}0)(a_1 : A \text{ bi}1)(aa : \text{BridgeP}_A a_0 a_1) \rightarrow \text{BridgeP}_{x.Bx}(aa x)(f_0 a_0)(f_1 a_1)}{\Gamma \vdash \text{extent } \{A\} \{B\} f_0 f_1 fr : (@x : BI)(a : Ax) \rightarrow Bx a} \text{EXT} \\
\\
\frac{\text{premises of EXT} \quad \Gamma \vdash a_\varepsilon : A \text{ bi}\varepsilon}{\Gamma \vdash \text{extent } f_0 f_1 fr \text{ bi}\varepsilon a_\varepsilon \mapsto f_\varepsilon a_\varepsilon} \text{EXT-}\partial \\
\\
\frac{\text{premises of EXT} \quad (x : BI) \in \Gamma \quad \Gamma \vdash a : Ax \quad \text{sfresh}(\text{red } a, x)}{\Gamma \vdash \text{extent } f_0 f_1 fr x a \mapsto fr(\langle x \rangle a \text{ bi}0)(\langle x \rangle a \text{ bi}1)(\langle x \rangle (\text{red } a)) x : Bx a} \text{EXT-}\beta \\
\\
\frac{\Gamma \vdash A_\varepsilon : \text{Type} \quad \Gamma \vdash R : A_0 \rightarrow A_1 \rightarrow \text{Type}}{\Gamma \vdash \text{Gel}_{A_0 A_1} R : (@x : BI) \rightarrow \text{Type}} \text{GEL-FORM} \quad \frac{}{\Gamma \vdash \text{Gel}_{A_0 A_1} R \text{ bi}\varepsilon \mapsto A_\varepsilon : \text{Type}} \text{GEL-}\partial \\
\\
\frac{\Gamma \vdash a_\varepsilon : A_\varepsilon \quad \Gamma \vdash ar : R a_0 a_1}{\Gamma \vdash \text{gel } \{A_0\} \{A_1\} \{R\} a_0 a_1 ar : (@x : BI) \rightarrow \text{Gel}_{A_0 A_1} R x} \text{GEL-INTRO} \quad \frac{}{\Gamma \vdash \text{gel } a_0 a_1 ar \text{ bi}\varepsilon \mapsto a_\varepsilon : A_\varepsilon} \text{GEL-}\partial \\
\\
\frac{\Gamma \vdash Q : (@x : BI) \rightarrow \text{Gel}_{A_0 A_1} R x}{\Gamma \vdash \text{ungel } \{A_0\} \{A_1\} \{R\} Q : R(Q \text{ bi}0)(Q \text{ bi}1)} \text{GEL-ELIM} \quad \frac{\Gamma \vdash a_\varepsilon : A_\varepsilon \quad \Gamma \vdash ar : R a_0 a_1}{\Gamma \vdash \text{ungel } (\lambda x. \text{gel } a_0 a_1 ar x) \mapsto ar : R a_0 a_1} \text{GEL-}\beta \\
\\
\frac{\Gamma \vdash g_\varepsilon : \text{Gel}_{A_0 A_1} R x \quad (x : BI) \in \Gamma \quad \text{sfresh}(\text{red } g_\varepsilon, x) \quad \Gamma \vdash \langle x \rangle g_0 \text{ bi}\varepsilon = \langle x \rangle g_1 \text{ bi}\varepsilon \quad \Gamma \vdash \text{ungel } (\langle x \rangle (\text{red } g_0)) = \text{ungel } (\langle x \rangle (\text{red } g_1)) : R(\langle x \rangle g_0 \text{ bi}0)(\langle x \rangle g_0 \text{ bi}1)}{\Gamma \vdash g_0 = g_1 : \text{Gel}_{A_0 A_1} R x} \text{GEL-}\eta
\end{array}$$

Fig. 1. Rules of the `BridgeP`, `extent` and `Gel` primitives of Agda --bridges. Some premises are omitted.

In order to validate the left-to-right direction of the above principle, Agda --bridges postulates the existence of the `extent` primitive (see also the `EXT` rule of Fig. 1). Note that the type of `extent` rather ends with a type of affine functions. This difference is essentially cosmetic thanks to the `EXT- ∂` rule of `extent`.

`extent` : $\forall \{A : (@\text{tick } x : BI) \rightarrow \text{Type}\} \{B : (@\text{tick } x : BI) (a : Ax) \rightarrow \text{Type}\}$
 $(f_0 : (a_0 : A \text{ bi}0) \rightarrow B \text{ bi}0 a_0) (f_1 : (a_1 : A \text{ bi}1) \rightarrow B \text{ bi}1 a_1)$
 $(fr : (a_0 : A \text{ bi}0) (a_1 : A \text{ bi}1) (aa : \text{BridgeP } A a_0 a_1) \rightarrow \text{BridgeP } (\lambda x \rightarrow Bx (aa x)) (f_0 a_0) (f_1 a_1))$
 $(@\text{tick } x : BI) (a : Ax) \rightarrow Bx a$

It remains to upgrade this map into an actual equivalence `extentEquiv`. Validating the right-to-left direction of `extentEquiv` is possible without assuming further primitives. Indeed, assuming a bridge `ff` : `Bridge`_{(a:A)→B} $f_0 f_1$, we can easily build a function of the appropriate type $\lambda a_0 a_1 aa \rightarrow \lambda x. ff x (aa x)$. Furthermore, one of the inverse conditions can be obtained using a “propositional η -rule” theorem proved for `extent`. The other inverse condition relies on the β -rule of `extent` (see `EXT- β` in Fig. 1). The exact proof can be consulted in our accompanying library, or in the CH paper.

The β -rule of **extent** is not a standard type theoretic rule, as it works by *capturing* (see **CAP** rule) the bridge variable argument $x : \mathbf{BI}$ of **extent** in its principal argument $a : Ax$. Capturing can only be performed if a specific freshness constraint (semi-freshness, denoted $\text{sfresh}(a, x)$) is satisfied and thus **EXT- β** only fires in that case. We now explain what semi-freshness is and how capturing is implemented. Note that several other parametricity primitives of Agda --bridges like **Gel** and the Kan operations make use of capturing in their equations.

Definition 2.2 (Semi-freshness constraint). Given a context $\Gamma = \Gamma_1, (x : \mathbf{BI}), \Gamma_2$ and a term $\Gamma \vdash a$, the constraint $\text{sfresh}(a, x)$ is satisfied if for every free variable v of a one of the following holds:

- $v \in \Gamma_1$ **or** $v = x$,
- $(v : \mathbf{I}) \in \Gamma_2$ **or**³ $(v : \mathbf{BI}) \in \Gamma_2$. We define Υ_2 as the variables v satisfying this clause.

If $\Gamma_1, (x : \mathbf{BI}), \Gamma_2 \vdash a : Ax$ and $\text{sfresh}(a, x)$ then a is in fact a weakening $a = a'[\pi]$ of a term $\Gamma_1, (x : \mathbf{BI}), \Upsilon_2 \vdash a' : A'x$ along $\pi : \Gamma_1, (x : \mathbf{BI}), \Gamma_2 \rightarrow \Gamma_1, (x : \mathbf{BI}), \Upsilon_2$ ⁴. Moreover, there is a well-typed substitution $\rho : \Gamma_1, (x : \mathbf{BI}), \Gamma_2, (y : \mathbf{BI}) \rightarrow \Gamma_1, (x : \mathbf{BI}), \Upsilon_2$ defined by $\rho = (\text{id}_{\Gamma_1}, y/x, \text{id}_{\Upsilon_2})$, so that we have $\Gamma_1, (x : \mathbf{BI}), \Gamma_2, (y : \mathbf{BI}) \vdash a'[\rho] : A'[\rho]y$. Agda --bridges will shortcut this by constructing a potentially ill-typed substitution $\sigma = (\text{id}_{\Gamma_1}, y/x, \text{id}_{\Gamma_2}) : \Gamma_1, (x : \mathbf{BI}), \Gamma_2, (y : \mathbf{BI}) \rightarrow \Gamma_1, (x : \mathbf{BI}), \Gamma_2$ which has the property that $\pi \circ \sigma = \rho$, and therefore $\Gamma_1, (x : \mathbf{BI}), \Gamma_2, (y : \mathbf{BI}) \vdash a[\sigma] = a'[\rho] : A[\sigma]y$ is well-typed. By freshness w.r.t. x , $A[\sigma] = A$ and by lambda abstraction we obtain $\Gamma \vdash \lambda y. a[\sigma] : (@y : \mathbf{BI}) \rightarrow Ay$. The latter is the definition of capturing $\Gamma \vdash \langle x \rangle a : (@y : \mathbf{BI}) \rightarrow Ay$.

2.5 Gel Types

The univalence theorem stated in Section 2.1.3 posits that paths at the universe $AA : A_0 \equiv A_1$ uniquely correspond to type equivalences $A_0 \simeq A_1$. Analogously, the *relativity* theorem asserts that *bridges* at the universe uniquely correspond to *relations*.

$$\text{relativity} : (A_0 \rightarrow A_1 \rightarrow \text{Type}) \simeq \text{Bridge}_{\text{Type}} A_0 A_1$$

Similar to **extent**, Agda --bridges postulates the existence of **Gel** in order to validate the left-to-right direction of **relativity**. Thus **Gel** has the following type (see also **GEL-FORM** in Fig. 1).

$$\text{Gel} : \forall (A_0 A_1 : \text{Type}) (R : A_0 \rightarrow A_1 \rightarrow \text{Type}) (@\text{tick } x : \mathbf{BI}) \rightarrow \text{Type}$$

The type of **Gel** merely ends with a type of affine functions $(@x : \mathbf{BI}) \rightarrow \text{Type}$ which can be turned into a bridge type thanks to the **GEL- ∂** rule.

To upgrade this map into the **relativity** equivalence, we first need an inverse candidate. From a bridge between two types $AA : \text{Bridge}_{\text{Type}} A_0 A_1$ we obtain the following relation $\lambda a_0 a_1 \rightarrow \text{BridgeP}_{x.AAx} a_0 a_1 : A_0 \rightarrow A_1 \rightarrow \text{Type}$. This relation between the types A_0 and A_1 holds for a_0, a_1 exactly when there is a dependent bridge over AA between them. We also need to provide two inverse conditions.

2.5.1 The First Inverse Condition. Regarding the first condition, using function extensionality one has to show $\forall (R : A_0 \rightarrow A_1 \rightarrow \text{Type}) (a_0 : A_0) (a_1 : A_1) \rightarrow \text{BridgeP}_{x.AAx} a_0 a_1 \rightarrow \text{BridgeP}_{x.AAx} a_0 a_1 \equiv R a_0 a_1$. By univalence, it can be reduced to proving an equivalence $\text{BridgeP}_{x.AAx} a_0 a_1 R x a_0 a_1 \simeq R a_0 a_1$. One could call this equivalence a (dependent) relational extensionality principle for **Gel**. Constructing both directions of this particular equivalence is precisely the role played by the introduction and elimination primitives of **Gel**, called **gel** and **ungel**, respectively (see **GEL-INTRO**, **GEL-ELIM**). The fact that **gel** and **ungel** are mutual inverses is in turn guaranteed by the equations governing those primitives, called the η -rule and the β -rule of **Gel** (see **GEL- β** , **GEL- η**). Note that the η -rule of **Gel**

⁴Because Ax is a well-typed bridge application, A is fresh w.r.t. x and therefore $A = A'[\pi]$.

makes use of capturing, as was the case for $\text{EXT-}\beta$. This means in particular that a semi-freshness constraint is checked when comparing two inhabitants of `Gel`.

2.5.2 The Second Inverse Condition. Regarding the second inverse condition needed to prove `relativity`, one has to show $\forall (AA : \text{Bridge } A_0 A_1) \rightarrow \text{Gel } A_0 A_1 (\lambda a_0 a_1. \text{BridgeP}_{x.AA x} a_0 a_1) \equiv AA$. Formally proving this lemma in Agda --bridges was far from trivial (see our accompanying library). In their paper and technical report, [Cavallo and Harper \[2019, 2021\]](#) sketch a proof of this lemma based on a relational extensionality principle for *equivalences* whose lengthy proof they also sketch. The principle is a characterization of the type of heterogeneous bridges between equivalences and its formulation is somewhat comparable to `extentEquiv`. We merely indicate here that our formal proof involves the pasting of several 2-dimensional paths in `Type`, whose construction relies on the path interval operations \sim, \wedge, \vee provided by Agda --cubical.

2.6 Other Relational Extensionality Principles

The primitives `extent` and `Gel, gel, ungel` we have implemented grant relational extensionality principles for the Π type former and for the universe `Type`, respectively. It turns out that their addition alone ensures the validity of similar principles for the other primitive type formers of the theory. For instance, as is the case for paths, a bridge at $\Sigma[a \in A] B a$ between pairs $(a_0, b_0), (a_1, b_1)$ can equivalently be regarded as a bridge in the base $aa : \text{Bridge}_A a_0 a_1$ together with a heterogeneous bridge over it $bb : \text{BridgeP}_{x.B(aa x)} b_0 b_1$. There also exist relational extensionality principles for the path type \equiv and the `Bridge` type itself. Essentially, those two principles reflect the fact that it is always possible to swap the order of bridge and path variables in the context.

Relational extensionality principles for specific inductive data types can also be stated and proved in Agda --bridges. For instance, in their paper CH prove such a principle for the type of booleans `Bool`. The principle expresses that `Bool` is *bridge-discrete*, that is, the only bridges in `Bool` are the ones corresponding to its paths: $b_0 \equiv_{\text{Bool}} b_1 \simeq (\text{Bridge}_{\text{Bool}} b_0 b_1)$. Since `Bool` has no non-reflexivity paths (i.e., is an h-set in HoTT parlance), there are only two bridges in `Bool`: the reflexivity bridges at `true` and `false`. We have adapted the argument of CH to prove in Agda --bridges a similar principle for the `List` parametrized data type. More precisely, it is a *dependent* extensionality principle as it characterizes the type of *heterogeneous* bridges $\text{BridgeP}_{x.\text{List}(AA x)} as_0 as_1$ between two lists $(as_0 : \text{List } A_0), (as_1 : \text{List } A_1)$ where $AA : \text{Bridge } A_0 A_1$.

```
ListRel : ∀ {A0 A1 : Type} (R : A0 → A1 → Type) → List A0 → List A1 → Type
ListRel R [] [] = Unit
ListRel R [] ( _ :: _ ) = ⊥
ListRel R ( _ :: _ ) [] = ⊥
ListRel R (a0 :: as0) (a1 :: as1) = R a0 a1 × ListRel R as0 as1
```

```
ListvsBridgeP : ∀ {A0 A1 : Type} (AA : Bridge A0 A1) (as0 : List A0) (as1 : List A1) →
  ListRel (BridgeP (λ x → AA x)) as0 as1 ≃ BridgeP (λ x → List (AA x)) as0 as1
```

The principle essentially expresses that a bridge between lists is a list of bridges. Indeed, `ListRel R` is an inductively defined relation between the types `List A0` and `List A1` which holds for as_0, as_1 exactly when the latter lists have the same size and exhibit R -related values at each index.

2.7 Low-Level Parametricity

We are now ready to prove theorems for free from first principles in Agda --bridges. It is customary in type theories with internal parametricity (incl. CH) to prove free theorems by directly appealing


```

lowChurchBool : (∀ (X : Type) → X → X → X) ≈ Bool
lowChurchBool = isoToEquiv (iso chToBool boolToCh (λ { true → refl ; false → refl })
  λ k → funExt λ A → funExt λ t → funExt λ f → param-prf k A t f)
where
  boolToCh : Bool → (∀ (X : Type) → X → X → X)
  boolToCh true X xt xf = xt
  boolToCh false X xt xf = xf

  chToBool : (∀ (X : Type) → X → X → X) → Bool
  chToBool k = k Bool true false

module CH-inverse-cond (k : ∀ (X : Type) → X → X → X) (A : Type) (t f : A) where
  R : Bool → A → Type
  R = λ b a → (boolToCh b A t f) ≡ a
  k-Gelx : (@tick x : Bl) → Gel Bool A R x → Gel Bool A R x → Gel Bool A R x
  k-Gelx x = k (Gel Bool A R x)
  k-Gelx-gel-gel : (@tick x : Bl) → Gel Bool A R x
  k-Gelx-gel-gel x = k-Gelx (gel true t (refl) x) ((gel false f (refl) x))
  asBdg : BridgeP (λ x → Gel Bool A R x) (k Bool true false) (k A t f)
  asBdg x = k-Gelx-gel-gel x
  param-prf : R (k Bool true false) (k A t f)
  param-prf = ungel {R = R} λ x → asBdg x
open CH-inverse-cond

```

Fig. 2. A low-level proof of a free theorem.

to the available low-level parametricity primitives. This is unsurprising since, after all, those primitives have been added precisely for that purpose.

Such a low-level proof of a free theorem in Agda --bridges appears in Fig. 2. The `lowChurchBool` theorem asserts that `Bool` admits a Church encoding, i.e., that this equivalence holds: $(X : \text{Type}) \rightarrow X \rightarrow X \rightarrow X \simeq \text{Bool}$. It is a faithful reproduction of the proof appearing in [Cavallo and Harper 2021]. We provide a high-level description of the proof and refer to the latter for more detailed explanations. To build an equivalence, it is sufficient to provide two maps and two inverse conditions. This is the content of the `isoToEquiv` lemma. The two candidate inverses are defined in a `where` block below and are called `boolToCh` and `chToBool`. The first inverse condition can simply be proven by induction. Using function extensionality, the second inverse condition asks that for every $k : (X : \text{Type}) \rightarrow X \rightarrow X \rightarrow X$, this equality holds $\text{boolToCh}(\text{chToBool } k) A t f \equiv k A t f$. Note that the universal quantification on k appears inside the system (with an Agda Π -type). The logic of Agda --cubical alone would not be sufficient to warrant this result (see Section 1). Therefore the internal parametricity of Agda --bridges must be used. All calls to parametricity primitives are isolated in a separate `module` called `CH-inverse-cond`. The last lemma of this module `param-prf` implies the second inverse condition.

We observe that such low-level proofs suffer from several defects. First, the user of Agda --bridges wanting to reproduce this style of proofs must have a good familiarity with the parametricity primitives provided by Agda --bridges, including their inner workings. But these primitives use advanced or non-standard type-theoretic notions like freshness and capturing. This makes for hard and non user-friendly proofs. Second these proofs lack compositionality. Indeed, we have proved in Fig. 2 a free theorem `param-prf` about polymorphic programs of type $T = (X : \text{Type}) \rightarrow X \rightarrow X \rightarrow X$. We expect other free theorems to hold at this type and it is unclear at first glance if `param-prf` could be reused to achieve that. In fact we even would like to be able to reuse the

latter parametricity proof to shortcut proofs of free theorems at types having $- \rightarrow - \rightarrow -$ as a subexpression. Lastly, these proofs are typically long even for seemingly simple examples of free theorems. Furthermore their complexity quickly gets intractable when the size of the target type T grows (there are several reasons for this, discussed in the next section).

All in all, these drawbacks motivated us to develop in Agda `--bridges` a library providing user-friendly, compositional and short proofs of free theorems. This is the content of Section 3.

3 THE OBSERVATIONAL PARAMETRICITY OF AGDA BRIDGES

As explained in Section 2.7, low-level proofs of internal free theorems are unsatisfactory in several respects. We improve these low-level proofs in two steps.

Our first improvement stems from the observation that, in order to make use of internal parametricity, it is always sufficient to prove appropriate relational extensionality principles. More precisely, we argue that obtaining internal free theorems for an Agda `--bridges` program $p : (\gamma : \Gamma) \rightarrow T \gamma$ can be reduced to providing (dependent) relational extensionality principles for p 's domain and codomain, so characterizations of their (dependent) bridge types as types of actual logical relations: an equivalence $\eta^\Gamma : \dots \simeq \text{Bridge}_\Gamma \gamma_0 \gamma_1$ and an equivalence $\eta^T : \dots \simeq \text{BridgeP}_{x.T(\gamma\gamma x)} (t_0 : T \gamma_0)(t_1 : T \gamma_1)$ where $\gamma\gamma : \text{Bridge}_\Gamma \gamma_0 \gamma_1$. This is explained in Section 3.1 and illustrated on an example in Section 3.1.4.

We call this informal sufficient condition for obtaining free theorems, which asks that all (dependent) types are equipped with a characterization of their `Bridge/BridgeP` type as logical relations, the *structure relatedness principle* (SRP). The SRP is precisely stated in Section 3.1. The reason behind this name is that there exists an analogous principle asking instead that all (dependent) types feature a characterization of their `≡/PathP` types as types of isomorphisms. The latter principle is known (to varying degrees of generality, see Section 6.3) as the *structure identity principle* (SIP) in HoTT/UF.

The second improvement we make compared to low-level proofs stems from the observation that proving the SRP or the SIP “by hand” at a given type can quickly get intractable, as explained in Section 3.2. To remedy this situation we introduce in Section 3.3 a shallowly embedded domain-specific language (DSL) implemented as an Agda `--bridges` library that allow the user to (1) show the SRP at a type T by merely writing their type T in the DSL (using the rules in Section 3.3.1) and (2) derive free theorems for T in a straightforward manner (see the `param` theorem of Section 3.3.2). We call our DSL *relational observational type theory* (ROTT). By contrast with low-level proofs, ROTT provides abstractions to write user-friendly, modular and concise proofs.

3.1 The SRP and Bare Parametricity

The first improvement we make for better internal parametricity proofs is to systematically factor proofs of free theorems into two simpler statements: the structure relatedness principle (SRP) on one side, and bare parametricity on the other. We first explain these principles and then illustrate their use by deriving the global free theorem (4) of Section 1 in Agda `--bridges`.

3.1.1 Bare Parametricity. Bare parametricity is simply the fact that all programs defined in Agda `--bridges` have a canonical action on bridges. It can be summarized as the following `bare-param` program:

```

bare-param : ∀ {Γ : Type} {T : Γ → Type} (p : ∀ γ → T γ) (γ₀ γ₁ : Γ)
  (γγ : Bridge γ₀ γ₁) → BridgeP (λ x → T (γγ x)) (p γ₀) (p γ₁)
bare-param p γ₀ γ₁ γγ = λ x → p (γγ x)

```

3.1.2 The SRP, Relativistic Reflexive Graphs and the SIP. The structure relatedness principle (SRP) is the following metatheoretical principle:

```

record RRGraph : Type1 where
  field
    cr : Type
    lrel : cr → cr → Type
    requ : ∀ (a b : cr) → lrel a b ≈ Bridge a b
open RRGraph public

record DispRRG (Γ : RRGraph) : Type1 where
  field
    dcr : Γ.cr → Type
    dlrel : ∀ (γ0 γ1 : Γ.cr) (γr : Γ.lrel γ0 γ1) (a0 : dcr γ0) (a1 : dcr γ1) → Type
    drequ : (γ0 γ1 : Γ.cr) (γr : Γ.lrel γ0 γ1) (γγ : Bridge γ0 γ1) (γprf : γr [ (Γ.requ γ0 γ1) ] γγ)
      (a0 : dcr γ0) (a1 : dcr γ1) → dlrel γ0 γ1 γr a0 a1 ≈ BridgeP (λ x → dcr (γγ x)) a0 a1
open DispRRG public

→Form : ∀ {Γ : RRGraph} (A B : DispRRG Γ) → DispRRG Γ
→Form A B.dcr γ = (A.dcr γ → B.dcr γ)
→Form A B.dlrel γ0 γ1 γr f0 f1 = ∀ a0 a1 → (ar : A.dlrel γ0 γ1 γr a0 a1) → B.dlrel γ0 γ1 γr (f0 a0) (f1 a1)
→Form A B.drequ γ0 γ1 γr γγ γprf f0 f1 = flip compEquiv extentEquiv --under the hood: extent
  ((equivΠCod λ a0 → equivΠCod λ a1 →
    equivIT' (A.drequ γ0 γ1 γr γγ γprf a0 a1) λ {ar} {aa} aprf →
    B.drequ γ0 γ1 γr γγ γprf (f0 a0) (f1 a1)))

```

Fig. 3. (Displayed) relativistic reflexive graphs in Agda --bridges, and their \rightarrow_{FM} semantic rule.

Structure Relatedness Principle (SRP): For each type $\Gamma : \text{Type}$ of Agda --bridges, (resp. type family $T : \Gamma \rightarrow \text{Type}$) the **Bridge** type of Γ (resp. the **BridgeP** type of T) can be characterized as a type of logical relations (resp. heterogeneous logical relations).

In other words, the SRP conveys the idea that bridges act as logical relations *at all types*^{*}. For instance the SRP holds at **Type** thanks to the **relativity** theorem of Section 2.5, and the SRP holds at function types $(a : A) \rightarrow B$ thanks to the **extentEquiv** theorem of Section 2.4.

Contrary to bare parametricity, the SRP can not be stated as an object-level theorem of Agda --bridges. The reason is that types of logical relations are ad hoc: there is a priori no Agda --bridges function **Lrel** : **Type** \rightarrow **Type** acting as an internal parametricity translation, computing for each Γ its type of logical relations **Lrel** Γ . Indeed parametricity translations are defined by induction on the syntax, so using a type-casing operation. But having a first-class type-casing operator on types would contradict the internal parametricity of Agda --bridges! This means that the quantification (*) must be stated metatheoretically.

Since the SRP can not be obtained as a theorem, we package it as a definition on types (or type families). For reasons explained hereafter, types that satisfy the SRP are called *relativistic reflexive graphs* (RRG). Moreover type families $T : \Gamma \rightarrow \text{Type}$ satisfying the SRP are called *displayed relativistic reflexive graphs* (dRRG). In what follows RRGs and dRRGs are defined in plain English. The corresponding formal Agda --bridges definitions of these structures are provided in Fig. 3.

Recall that we use a postfix r notation to denote (potentially heterogeneous) logical relations ar and a double-letter notation aa to denote (potentially dependent) bridges. Additionally, note that we use the notation $a_0 [e] a_1$ to signify that $e.\text{fst } a_0 \equiv_{A_1} a_1$ where $e : A_0 \approx A_1$ and $a_0 : A_0, a_1 : A_1$.

Definition 3.1 (Relativistic reflexive graph (RRG)). A relativistic reflexive graph is a type $\Gamma : \text{Type}$ which, for all elements $\gamma_0, \gamma_1 : \Gamma$, is equipped with a type denoted $\Gamma\{\gamma_0, \gamma_1\} : \text{Type}$ and an equivalence denoted $\eta^\Gamma : \Gamma\{\gamma_0, \gamma_1\} \approx \text{Bridge}_\Gamma \gamma_0 \gamma_1$.

Members of $\Gamma\{\gamma_0, \gamma_1\}$ are called logical relations at Γ between γ_0, γ_1 . The η equivalence is called the relativistic equivalence of Γ . When the context is clear, we allow ourselves to refer to the RRG given by the triple $(\Gamma, \lambda\gamma_0 \gamma_1 \rightarrow \Gamma\{\gamma_0, \gamma_1\}, \lambda\gamma_0 \gamma_1 \rightarrow \eta^\Gamma)$ simply as Γ . Next, the notion of displayed relativistic reflexive graph (dRRG) is an indexed version of the above notion. The definition is not straightforward but exactly encodes what we want: types that carry characterizations of their **BridgeP** types. To help the reader parse the definition, the indexed counterpart of each RRG operation is written in **bold** font.

Definition 3.2 (Displayed relativistic reflexive graph over $-$ (dRRG over $-$)). Let Γ be a RRG. A displayed relativistic reflexive graph T over Γ is

- For every $\gamma : \Gamma$ **a type denoted** T_γ with...
- For every $\gamma_0 \gamma_1$ ($\gamma r : \Gamma\{\gamma_0, \gamma_1\}$) ($t_0 : T \gamma_0$) ($t_1 : T \gamma_1$) **a type denoted** $T\{t_0, t_1\}_{\gamma r}$ with ...
- For every $(\gamma r : \Gamma\{\gamma_0, \gamma_1\})$ and $(\gamma \gamma : \mathbf{Bridge}_\Gamma \gamma_0 \gamma_1)$ such that $\gamma r [\eta^\Gamma] \gamma \gamma$ and for every $(t_0 : T \gamma_0), (t_1 : T \gamma_1)$, **an equivalence denoted** $\eta^T : T\{t_0, t_1\}_{\gamma r} \simeq \mathbf{BridgeP}_{x.T(\gamma \gamma x)} t_0 t_1$.

Members of $T\{t_0, t_1\}_{\gamma r}$ are called heterogeneous logical relations at T between t_0, t_1 over the logical relation γr . If the triple $(T, \lambda... \rightarrow T\{t_0, t_1\}_{\gamma r}, \lambda... \rightarrow \eta^T)$ is a dRRG over Γ , we allow ourselves to refer to it simply as T . In that case, we also adopt the suggestive notation $(\gamma : \Gamma) \models T_\gamma$ dRRG by analogy with the *type* judgment of type theory $(x_0 : A_0, \dots, x_n : A_n \vdash T \text{ type})$.

We motivate our terminology of RRGs and dRRGs for types that satisfy the SRP. First, types Γ that have the SRP, i.e., are equipped with types $\Gamma\{\gamma_0, \gamma_1\}$ and an equivalence $\eta : \Gamma\{\gamma_0, \gamma_1\} \simeq \mathbf{Bridge}_\Gamma \gamma_0 \gamma_1$, are reflexive graphs. Indeed we can regard the logical relations of Γ as its edges and we can pick $\eta^{-1}(\lambda x. \gamma) : \Gamma\{\gamma, \gamma\}$ for its reflexivity edges. Second, the archetypal type satisfying the SRP is **Type**, thanks to **relativity**. Hence types that satisfy the SRP are called *relativistic reflexive graphs*. Following Ahrens and Lumsdaine's [2019] convention for naming dependent mathematical objects, we call the dependent version of a RRG a displayed RRG.

By exact analogy with the SRP, the structure identity principle (SIP; see Section 6.3 for references) is a metatheoretical statement asking that, for each type Γ (resp. $T : \Gamma \rightarrow T$) the \equiv type of Γ (c.f. bridge type; resp. **PathP** type of T) can be characterized as a type of *isomorphisms* (c.f. logical relations). That is to say, the SIP conveys the idea that paths act as isomorphisms at all types. For instance the SIP holds at **Type** and at function types $(a : A) \rightarrow B$ thanks to the **univalence** and **funextEquiv** theorems of Section 2.1.3. Types satisfying the SIP are most commonly known as univalent groupoids, or setoids in the literature (c.f. relativistic reflexive graphs).

3.1.3 Examples of RRGs and dRRGs. We give two examples of RRGs and one example of dRRG. The upshot is that each (dependent) relational extensionality principle that we can prove (see Section 2.6) for a potentially composite type, can be repackaged as a (d)RRG.

First, we define the composite type of premonoids **PreMon** $= \Sigma[M \in \mathbf{Type}] M \times (M \rightarrow M \rightarrow M)$. This type can be turned into an RRG. Indeed we can pick $\Gamma = \mathbf{PreMon}$ and define $\mathbf{PreMon}\{M_0, M_1\}$ as the type of (actual) logical relations of premonoids between M_0, M_1 . A logical relation of premonoids is a relation between M_0, M_1 with a proof that the neutral elements of M_0, M_1 are related, and a proof that the binary functions are pointwise related. By combining the principles for $\Sigma, \times, \mathbf{Type}, \rightarrow$ appearing in Section 2.6, one can prove that the type $\mathbf{PreMon}\{M_0, M_1\}$ is equivalent to $\mathbf{Bridge}_{\mathbf{PreMon}} M_0 M_1$ which makes $\Gamma = \mathbf{PreMon}$ an RRG. Second, the type **Bool** \rightarrow **Bool** can be turned into an RRG. We can indeed pick $(\mathbf{Bool} \rightarrow \mathbf{Bool})\{f_0, f_1\} = (\forall b_0 b_1 \rightarrow f_0 b_0 \equiv_{\mathbf{Bool}} f_1 b_1)$ and use the appropriate relational extensionality principles, including the one for **Bool** mentioned in (2.6) to build the expected equivalence. Our third example regards the **List** type former. Recall that the triple $(\mathbf{Type}, \lambda A_0 A_1 \rightarrow A_0 \rightarrow A_1 \rightarrow \mathbf{Type}, \mathbf{relativity})$ turns **Type** into a RRG. We assert that the type

family $\lambda X \rightarrow \text{List } X$ is a displayed RRG over the **Type** RRG, i.e., $(X : \text{Type}) \models \text{List } X \text{ dRRG}$. The reason is that we can pick $(\lambda X \rightarrow \text{List } X) \{as_0, as_1\}_R = \text{ListRel } R \text{ } as_0 \text{ } as_1$, the inductive characterization of $\text{BridgeP}_{x.\text{List}(AA\ x)}$ discussed in Section 2.6.

3.1.4 SRP + Bare Parametricity. Next, we illustrate how the SRP and bare parametricity can be used to improve proofs of internal free theorems. The main idea is that since (1) bare parametricity tells us that programs act on bridges and (2) the SRP guarantees that bridges uniquely correspond to logical relations, we expect that (3) programs act on logical relations as well. And all free theorems are consequences of the latter fact. We derive the global free theorem (4) of Section 1 in Agda –bridges, or rather a slightly reduced version of it for sparing space.

$$\begin{aligned} \text{fthm} : & \forall \{A_0 \ A_1 : \text{Type}\} (f : A_0 \rightarrow A_1) (as_0 : \text{List } A_0) \\ & (p : (X : \text{Type}) \rightarrow \text{List } X \rightarrow \text{List } X) \rightarrow \text{map } f (p \ A_0 \ as_0) \equiv p \ A_1 \ (\text{map } f \ as_0) \end{aligned}$$

The first step of our proof is to derive the SRP for the domain and codomain of p . In other words we must (1) equip **Type** with an RRG structure (we chose the one induced by **relativity**) and (2) equip the type family $\lambda X. \text{List } X \rightarrow \text{List } X$ with a dRRG structure over the **Type** RRG. This amounts to proving the following characterization of the **BridgeP** type of $\lambda X. \text{List } X \rightarrow \text{List } X$ where $A_0, A_1 : \text{Type}$, $R : A_0 \rightarrow A_1 \rightarrow \text{Type}$, $AA : \text{Bridge } A_0 \ A_1$, $\text{Aprf} : R \text{ [relativity] } AA$, and $q_\varepsilon : \text{List } A_\varepsilon \rightarrow \text{List } A_\varepsilon$ for $\varepsilon = 0, 1$. The proof uses the relational extensionality principles **extentEquiv**, **ListvsBridgeP** and the one for **Gel** (see Section 2.5.1, 2.6) as well as the fact that all type formers preserve equivalences. In the spirit of parametricity translations, an appropriate relational extensionality principle is used based on the head of the type former appearing in the **BridgeP** type at hand.

$$\begin{aligned} & \text{BridgeP}_{x.\text{List}(AA\ x) \rightarrow \text{List}(AA\ x)} \ q_0 \ q_1 \quad \simeq \\ & \forall as_0 \ as_1 \rightarrow \text{BridgeP}_{x.\text{List } AA\ x} \ as_0 \ as_1 \rightarrow \text{BridgeP}_{x.\text{List}(AA\ x)} \ (q_0 \ as_0) (q_1 \ as_1) \quad \simeq \\ & \forall as_0 \ as_1 \rightarrow (\text{ListRel } (\text{BridgeP}_{x.AA\ x}) \ as_0 \ as_1) \rightarrow (\text{ListRel } (\text{BridgeP}_{x.AA\ x}) \ (q_0 \ as_0) (q_1 \ as_1)) \simeq \\ & \forall as_0 \ as_1 \rightarrow (\text{ListRel } R \ as_0 \ as_1) \rightarrow (\text{ListRel } R \ (q_0 \ as_0) (q_1 \ as_1)) \end{aligned}$$

The second step of our proof is to “conjugate” **bare-param** with the SRP proof obligations we produced for **Type** and $\lambda X. \text{List } X \rightarrow \text{List } X$. Let A_0, A_1, f, as_0, p be as in **fthm**. We first convert the graph relation of f denoted $\text{Gr } f = \lambda a_0 \ a_1 \rightarrow f \ a_0 \equiv_{A_1} a_1$ into a bridge denoted $AA : \text{Bridge}_{\text{Type}} \ A_0 \ A_1$, using **relativity**. Then we apply bare parametricity to AA .

$$\text{bare-param } \{T = \lambda X \rightarrow \text{List } X \rightarrow \text{List } X\} \ p \ A_0 \ A_1 \ AA : \text{BridgeP}_{x.\text{List}(AA\ x) \rightarrow \text{List}(AA\ x)} \ (p \ A_0) (p \ A_1)$$

Finally we use the above dependent principle for $\lambda X. \text{List } X \rightarrow \text{List } X$ and obtain a proof $\text{pf} : \forall as_0 \ as_1 \rightarrow (\text{ListRel } (\text{Gr } f) \ as_0 \ as_1) \rightarrow (\text{ListRel } (\text{Gr } f) \ (p \ A_0 \ as_0) (p \ A_1 \ as_1))$. By a simple list induction we see that $\text{ListRel } (\text{Gr } f)$ is equal to the relation $\lambda as_0 \ as_1 \rightarrow (\text{map } f) \ as_0 \equiv as_1$. Thus $\text{pf } as_0 \ (\text{map } f \ as_0) \ \text{refl}$ grants the free theorem **fthm**.

The technique of factoring free theorems into SRP proof obligations and a call to **bare-param** is an improvement compared to low-level parametricity proofs like the one presented in Section 2.7. The proofs are conceptually easier. Moreover they allow for more compositionality. Indeed, contrary to Section 2.7 we are now in position to reuse the SRP proof obligations obtained for **Type** and $\lambda X. \text{List } X \rightarrow \text{List } X$ to derive (1) other free theorems for programs $p : (X : \text{Type}) \rightarrow \text{List } X \rightarrow \text{List } X$ and even (2) shorter proofs of free theorems for a composite type having **List** – \rightarrow **List** – as a subexpression. However it turns out that proving SRP obligations “by hand” like in the above is challenging, in fact strictly more challenging than SIP obligations, as explained in the next subsection.

3.2 Obstructions to the SRP and SIP

In this subsection we argue that proving the SIP, or worse, the SRP at a given type can get tedious. We begin with an example of SIP and SRP proofs for the type of pointed unary type operations $\text{PointedOp} = \Sigma[F \in \text{Type} \rightarrow \text{Type}] ((X : \text{Type}) \rightarrow X \rightarrow F X)$.

Example 3.3 (The SIP via extensionality principles). By repeatedly applying extensionality principles (see Section 2.1.3) we can characterize the meaning of a path between such pointed operations:

$$\begin{aligned}
 (F_0, f_0) &\equiv_{\text{PointedOp}} (F_1, f_1) \\
 &\simeq \Sigma[FF \in F_0 \equiv_{\text{Type} \rightarrow \text{Type}} F_1] \text{PathP}_{i.(X:\text{Type}) \rightarrow X \rightarrow FF i X} f_0 f_1 \\
 &\simeq \Sigma[FF' \in (X : \text{Type}) \rightarrow (F_0 X) \equiv_{\text{Type}} (F_1 X)] \text{PathP}_{i.(X:\text{Type}) \rightarrow X \rightarrow FF' X i} f_0 f_1 \\
 &\simeq \Sigma[e \in (X : \text{Type}) \rightarrow F_0 X \simeq F_1 X] \text{PathP}_{i.(X:\text{Type}) \rightarrow X \rightarrow \text{ua}(e X) i} f_0 f_1 \\
 &\simeq \Sigma[e \in (X : \text{Type}) \rightarrow F_0 X \simeq F_1 X] ((X : \text{Type}) \rightarrow (x : X) \rightarrow \text{PathP}_{i.\text{ua}(e X) i} (f_0 X x) (f_1 X x)) \\
 &\simeq \Sigma[e \in (X : \text{Type}) \rightarrow F_0 X \simeq F_1 X] ((X : \text{Type}) \rightarrow (x : X) \rightarrow f_0 X x [e X] f_1 X x)
 \end{aligned}$$

We conclude that a path between pointed operations (F_0, f_0) and (F_1, f_1) consists of a pointwise equivalence between F_0 and F_1 , compatible with the pointings f_0 and f_1 .

Example 3.4 (The SRP via relational extensionality principles). Set $\text{Rel } X_0 X_1 := X_0 \rightarrow X_1 \rightarrow \text{Type}$ and $\text{ra} := \text{relativity}$. We can characterize bridges at PointedOp by applying the principles discussed in Section 2.5.1, 2.6.

$$\begin{aligned}
 \text{Bridge}_{\text{PointedOp}} (F_0, f_0) (F_1, f_1) \\
 &\simeq \Sigma[FF : \text{Bridge}_{\text{Type} \rightarrow \text{Type}} F_0 F_1] \text{BridgeP}_{y.(X:\text{Type}) \rightarrow X \rightarrow FF y X} f_0 f_1 \\
 &\simeq \Sigma[FF' : (X_0 X_1 : \text{Type}) \rightarrow \text{Bridge}_{\text{Type}} X_0 X_1 \rightarrow \text{Bridge}_{\text{Type}} (F_0 X_0) (F_1 X_1)] \\
 &\quad (X_0 X_1 : \text{Type}) (XX : \text{Bridge}_{\text{Type}} X_0 X_1) (x_0 : X_0) (x_1 : X_1) \rightarrow \\
 &\quad \text{BridgeP}_{y.XX y x_0 x_1} \rightarrow \text{BridgeP}_{y.FF' X_0 X_1 XX y} (f_0 X_0 x_0) (f_1 X_1 x_1) \\
 &\simeq \Sigma[Fr : (X_0 X_1 : \text{Type}) \rightarrow \text{Rel } X_0 X_1 \rightarrow \text{Rel} (F_0 X_0) (F_1 X_1)] \\
 &\quad (X_0 X_1 : \text{Type}) (R : \text{Rel } X Y) (x_0 : X_0) (x_1 : X_1) \rightarrow \\
 &\quad \text{BridgeP}_{y.\text{ra } R y x_0 x_1} \rightarrow \text{BridgeP}_{y.\text{ra} (Fr X_0 X_1 R) y} (f_0 X_0 x_0) (f_1 X_1 x_1) \\
 &\simeq \Sigma[Fr : (X_0 X_1 : \text{Type}) \rightarrow \text{Rel } X_0 X_1 \rightarrow \text{Rel} (F_0 X_0) (F_1 X_1)] \\
 &\quad (X_0 X_1 : \text{Type}) (R : \text{Rel } X Y) (x_0 : X_0) (x_1 : X_1) \rightarrow \\
 &\quad (R x_0 x_1) \rightarrow Fr X_0 X_1 R (f_0 X_0 x_0) (f_1 X_1 x_1)
 \end{aligned}$$

We conclude that a bridge between pointed operations (F_0, f_0) and (F_1, f_1) consists of a relation transformer Fr between them such that the pointings f_0 and f_1 send related pairs to related pairs.

The examples above required rote work: we had to apply a series of extensionality principles which was entirely dictated by the formation of PointedOp . Hence, it is clear that proofs of the SIP and SRP at a type T scale at least with the complexity of T : for each type former F used to define T , an appropriate extensionality principle must be used to swap $\text{PathP}/\text{BridgeP}$ and F .

Moreover we argue that proving the SRP is in general strictly harder than proving the SIP, because of three obstructions which we list here.

The first obstruction is that the extentEquiv principle of Section 2.4 always produces an extra Bridge type when characterizing the type of bridges between two given functions. Both generated Bridge types must be further characterized to finish the SRP proof at hand. This is to compare with the funextEquiv principle of Section 2.1.3 which (if not in the fully dependent case) does not produce an extra PathP type.

The second and third obstructions relate to the fact that, in the SIP case, some tools are available to dismiss part of the rote work seen above. A first such tool is a theorem (see Rijke [2022] e.g.) that can be seen as a reformulation in cubical type theory of the J rule of equality types. Recall that a type A is contractible ($\text{isContr } A$) if it is equivalent to the unit type $\{*\}$.

THEOREM 3.5 (FUNDAMENTAL THEOREM OF IDENTITY TYPES). *Assume $A : \text{Type}$ and a relation $\text{Eq} : A \rightarrow A \rightarrow \text{Type}$. Suppose that Eq is reflexive, i.e., $\forall a \rightarrow \text{Eq } a \ a$. Additionally, suppose that $\forall a_0 \rightarrow \text{isContr } (\Sigma [a_1 \in A] \text{Eq } a_0 \ a_1)$. Then $\forall a_0 \ a_1 \rightarrow (\text{Eq } a_0 \ a_1) \simeq (a_0 \equiv_A a_1)$.*

This theorem might allow us to shortcut proofs as in Example 3.3, especially in the case of ordinary data types, by simply proving that the end result satisfies the above criteria. However, since bridges have no elimination principles like J , the theorem does not to our knowledge translate to the relational setting.

Similarly, the third obstruction is the lack of the following result, standard in HoTT but only available in case of the SIP. Recall that a type P is an h-proposition if any two inhabitants of P are equal, i.e., $\text{isProp } P = \forall p_0 \ p_1 \rightarrow p_0 \equiv_P p_1$. For instance the empty and unit types are h-propositions. The theorem guarantees the existence and unicity of heterogeneous paths between proofs of h-propositions.

THEOREM 3.6 (HETEROGENOUS EQUALITY OF PROOFS). *Let $P : I \rightarrow \text{Type}$ be a line of h-propositions, i.e., there is a map $\text{isp} : (i : I) \rightarrow \text{isProp } (P \ i)$. Suppose that $p_0 : P \ i_0$ and that $p_1 : P \ i_1$. Then $\text{isContr } (\text{Path}_{P, i_0 \ i_1} p_0 \ p_1)$ ⁵.*

Typically, this theorem is used while proving the SIP for structured types of the form $\Sigma [a \in A] (P \ a)$ where for all a , $P \ a$ is an h-proposition. For instance this is the reason why two equivalences $e_0, e_1 : A_0 \simeq A_1$ are equal if and only if their underlying functions are equal $e_0 \equiv e_1 \simeq (e_0.\text{fst} \equiv e_0.\text{fst})$. This is to compare with the relational extensionality principle for equivalences evoked in Section 2.5.2 which has the hardest proof amongst basic relational principles.

Because of these obstructions to SRP proofs, we have developed a domain-specific language (DSL) to obtain proofs of the SRP at a given type T by merely writing the type T in the DSL. Our DSL is explained in the next subsection.

3.3 ROTT

The second improvement we make compared to low-level proofs of free theorems is the introduction of *relational observational type theory* (ROTT): a domain-specific language (DSL) implemented as a library in Agda --bridges. It allows users to (1) show the SRP at a type T by merely writing their type T in the DSL and (2) derive free theorems for T in a straightforward manner.

Since ROTT has abstractions that compose *dependent types* T satisfying the SRP (i.e. dRRGs from Definition 3.2), it is itself a dependent type theory. To be more precise, ROTT is a type theory *shallowly embedded* in Agda --bridges. This means that ROTT is not defined as some kind of data type of expressions whose constructors would be inference rules used to write types T . Instead, ROTT is an Agda --bridges library (part of our accompanying library) comprised of a bunch of theorems called “semantic rules” explaining how to compose dRRGs. These Agda --bridges theorems are discussed in Section 3.3.1. As an example, we program in Fig. 3 the \rightarrow_{FM} semantic rule of ROTT directly in Agda --bridges.

Additionally, ROTT also features a `param` theorem letting the user straightforwardly deduce free theorems from appropriate (d)RRGs instances (obtained with ROTT, e.g.). We see `param` as one of the rules of ROTT. It is discussed in Section 3.3.2.

⁵For bridges, only $\text{isProp } (\text{BridgeP}_{x, P \ x} p_0 \ p_1)$ holds. By (2.5.1), $\text{BridgeP}_{x, \text{Gel } P_0 \ P_1} (\lambda - \dots \perp) \ x \ p_0 \ p_1 \simeq \perp$.

$$\begin{array}{c}
\frac{\Gamma \text{ RRG} \quad \Gamma \models T \text{ dRRG}}{\Gamma \# T \text{ RRG}} \text{CTX-EXT} \quad \frac{\Gamma \models A \text{ dRRG} \quad \Gamma \# A \models B \text{ dRRG}}{\Gamma \models \Pi \text{fm } A B \text{ dRRG}} \Pi \text{FM} \\
\frac{\Gamma \models A \text{ dRRG} \quad \Gamma \# A \models B \text{ dRRG}}{\Gamma \models \Sigma \text{fm } A B \text{ dRRG}} \Sigma \text{FM} \quad \frac{\Gamma \models f : \Pi \text{fm } A B \quad \Gamma \models a : A}{\Gamma \models \text{app } f a : B[a]} \text{APP} \\
\frac{\Gamma \text{ NRG}}{\Gamma \models \text{Tyfm } \text{dNRG}} \text{TYFM} \quad \frac{\Gamma \models A \text{ dRRG} \quad \Gamma \models a_0 : A \quad \Gamma \models a_1 : A}{\Gamma \models \equiv \text{fm } a_0 a_1 \text{ dRRG}} \equiv \text{FM}
\end{array}$$

Fig. 4. Some rules of ROTT. Some premises are not displayed.

3.3.1 The Standard Rules of ROTT. The key idea behind ROTT is to remark that RRGs Γ act as if they were *contexts* of a certain type theory, and displayed RRGs T over Γ act as if they were *types* of a type theory. In other words it seems like the type of all RRGs `RRGraph` somehow constitutes a model of type theory. For the expert reader this can be summarized as the fact that ROTT is a raw category-with-families (CwF) structure (with support for various type formers) on the category formed by relativistic reflexive graphs and their morphisms⁶. Here “raw” means that no equations between the CwF operations are required. There is no need to define what a (raw) CwF is since we only deal with one instance here and since in practice our proofs of the SRP sometimes combine manual reasoning with using the ROTT interface.

ROTT features standard (shallowly embedded!) type theoretic rules to combine RRGs and dRRGs, treating them as type-theoretic contexts and types in contexts, respectively. Some of those rules appear in Fig. 4 using traditional type-theoretic syntax. Recall that our notation for displayed RRGs over Γ is $\Gamma \models T \text{ dRRG}$. Each of those rules is an Agda `--bridges` program parametrized by the premises appearing in the rule. For example, the $\rightarrow \text{FM}$ rule of ROTT (a restricted version of ΠFM) is programmed directly in Agda `--bridges` in Fig. 3.

Let us explain how some of those rules are defined in Agda `--bridges`. The important idea here is that those rules compose the relational extensionality principles of their premises to yield composite types satisfying the SRP. This means that from the point of view of the user of ROTT, the SRP is proved by the system while types are being written with those rules. As can be expected, given Γ a RRG and T a displayed RRG over it, the context extension $\Gamma \# T$ is a RRG whose carrier is simply $\Sigma[\gamma \in \Gamma] T_\gamma$. Its type of logical relations $(\Gamma \# T)\{(\gamma_0, t_0), (\gamma_1, t_1)\}$ is moreover defined as $\Sigma[\gamma r \in \Gamma\{\gamma_0, \gamma_1\}] T\{t_0, t_1\}_{\gamma r}$. The fact that this type matches the corresponding bridge type `Bridge $\Gamma \# T$ (γ_0, t_0)(γ_1, t_1)`, can easily be deduced from the same fact for Γ and T . Regarding Πfm and Σfm , and supposing that Γ is empty, $\Pi \text{fm } A B$ and $\Sigma \text{fm } A B$ are RRGs⁷ with carriers $(a : A) \rightarrow B a$ and $\Sigma[a \in A](B a)$. In particular $((a : A) \rightarrow B a)\{f_0, f_1\}$ is defined as the type $\forall a_0 a_1 (ar : A\{a_0, a_1\}) \rightarrow B\{f_0 a_0, f_1 a_1\}_{ar}$ and this type can again be proven to match the corresponding `Bridge` type by composing the relativistic equivalences of A and B . The `Tyfm` rule equips `Type` with a dRRG structure in the expected way. For an empty Γ , the rule sets `Type` $\{A_0, A_1\} = A_0 \rightarrow A_1 \rightarrow \text{Type}$ and uses `relativity`. Finally ROTT also features a term judgment written $\Gamma \models a : A$ needed because arbitrary dependent types might mention terms in their definition. We do not state the definition of this judgment and merely indicate that $\Gamma \models a : A$ holds if $a : (\gamma : \Gamma) \rightarrow A_\gamma$ features a custom action on logical relations $\gamma r : \Gamma\{\gamma_0, \gamma_1\}$ which matches its canonical action on `Bridge Γ $\gamma_0 \gamma_1$` , i.e., `bare-param a $\gamma_0 \gamma_1$` .

3.3.2 Internal Observational Parametricity. We state the `param` semantic rule of ROTT. Once again we do so as if it was a syntactic type-theoretic rule even though `param` really is an Agda `--bridges` program parametrized by the premises appearing in the rule. The theorem states that, given an

⁶In our development those morphisms are called relativistic, or native relators but we will not need this notion here.

⁷There is a slight mismatch between dRRGs in empty contexts and RRGs, but we ignore it here.

RRG Γ and a displayed RRG T over it, all external dependent functions (i.e. all functions definable in Agda --bridges) from Γ to T respect logical relations.

$$\frac{\Gamma \text{ RRG} \quad \Gamma \models T \text{ dRRG} \quad p : (\gamma : \Gamma) \rightarrow T_\gamma \quad \gamma_0, \gamma_1 : \Gamma \quad \gamma r : \Gamma\{\gamma_0, \gamma_1\}}{\text{param } \Gamma \text{ } T \text{ } p \text{ } \gamma_0 \text{ } \gamma_1 \text{ } \gamma r : T\{p \text{ } \gamma_0, p \text{ } \gamma_1\}_{\gamma r}} \text{PARAM}$$

The proof of this theorem is elementary and proceeds in three steps, similar to the proof appearing in Section 3.1.4. We omit to write `.fst` to extract direct maps out of equivalences. First convert the logical relation γr into a bridge $\eta^\Gamma \gamma r : \text{Bridge}_\Gamma \gamma_0 \gamma_1$. Second use bare parametricity to obtain `bare-param` $p \text{ } \gamma_0 \text{ } \gamma_1 (\eta^\Gamma \gamma r)$ which has type $\text{BridgeP}_{x.T(\eta^\Gamma \gamma r x)} (p \text{ } \gamma_0) (p \text{ } \gamma_1)$. Third by definition we know that $\gamma r [\eta^\Gamma] (\eta^\Gamma \gamma r)$. Hence we may use the relativistic equivalence η^T of T to obtain $\eta^{T^{-1}}(\text{bare-param } p \text{ } \gamma_0 \text{ } \gamma_1 (\eta^\Gamma \gamma r))$ which has type $T\{p \text{ } \gamma_0, p \text{ } \gamma_1\}_{\gamma r}$. This concludes the proof and definition of `param`.

The `param` theorem draws inspiration from observational type theory. It can indeed be seen as a relational analogue of the *ap* inference rule (see e.g. [Altenkirch et al. 2022] and Section 6) which states that terms of observational type theory act on identifications or isomorphisms, that is, observational proofs of equality. For this reason we say that our internal `param`-etricity theorem is also observational. In the next section we use ROTT and its `param` rule to obtain modular and concise proofs of internal free theorems.

4 INTERNAL OBSERVATIONAL PARAMETRICITY APPLIED

In this section we obtain several free theorems as one-liner invocations of the `param` theorem of Section 3.3.2. This is done by first constructing appropriate (d)RRGs using the rules of ROTT. All of our examples can be consulted in our accompanying library.

4.1 Reproving `fthm` and `lowChurchBool`

We begin by recasting our proof of the global free theorem `fthm` from Section 3.1.4, this time using ROTT and its `param` rule. Let $p : (X : \text{Type}) \rightarrow \text{List } X \rightarrow \text{List } X$, let $f : A_0 \rightarrow A_1$ for A_0, A_1 two types and let $as : \text{List } A_0$. We want to apply `param` at program p and at (logical) relation $\text{Gr } f : A_0 \rightarrow A_1 \rightarrow \text{Type}$. In order to do so we must first provide an RRG structure for `Type`, the domain of p . This is achieved using the `Tyfm` rule of ROTT. Second, we must prove that $X : \text{Type} \models \text{List } X \rightarrow \text{List } X \text{ dRRG}$. By the `lfm` rule of ROTT (or rather \rightarrow_{FM} of Fig. 3) it is sufficient to prove (twice) that $X : \text{Type} \models \text{List } X \text{ dRRG}$. The latter displayed RRG appears as an example in Section 3.1.3. At this point all premises of `param` have been supplied.

Note how ROTT allows a significant improvement compared to proofs of the SRP “by hand” as shown Section 3.1.4. Instead of proving a lengthy equivalence chain, we simply have to write the following in Agda --bridges, using the `→Form` and `X≡ListX` programs implemented by the ROTT library (Agda identifiers can feature symbols, as in `X≡ListX`).

```
X≡ListX→ListX : DispRRG TypeRRG
X≡ListX→ListX = →Form X≡ListX X≡ListX
```

Looking at the conclusion of `param` we are about to obtain something of type $(\lambda X \rightarrow \text{List } X \rightarrow \text{List } X)\{p \text{ } A_0, p \text{ } A_1\}_{\text{Gr } f}$ and contrary to the `BridgeP` type former, the latter type *reduces* to the expected relational parametricity statement, i.e.:

$$\text{param TypeRRG X≡ListX→ListX } p \text{ } A_0 \text{ } A_1 (\text{Gr } f) : \\ \forall xs_0 xs_1 (xsr : \text{ListRel } (\text{Gr } f) xs_0 xs_1) \rightarrow \text{ListRel } (\text{Gr } f) (p \text{ } A_0 xs_0) (p \text{ } A_1 xs_1)$$

By applying this function to $xs_0 = as_0$, $xs_1 = \text{map } f \text{ } as_0$ and by remarking that $\text{ListRel } (\text{Gr } f) l_0 l_1$ is the same predicate than $\text{map } f \text{ } l_0 \equiv l_1$ we conclude the proof of `fthm`.

Similarly, we can reprove the `lowChurchBool` theorem of Section 2.7. The call to `param-prf` and the `CH-inverse-cond` module appearing in Fig. 2 are replaced by the following call to the `param` theorem and definition of an appropriate dRRG using ROTT:

```
param TypeRRG X=X→X→X k Bool A (λ b x → boolToCh b A t f ≡ x) true t refl false f refl
```

```
X=X→X→X : DispRRG TypeRRG
X=X→X→X = →Form (X=EIX) (→Form X=EIX X=EIX)
```

4.2 Church Encodings

Using ROTT and `param` we were able to prove a scheme of Church encodings for data types obtained out of a strictly positive functor, represented as a container [Abbott et al. 2005]. We first state the theorem and briefly explain its proof in Agda --bridges, and then discuss its hypotheses and significance in the more concrete case of the `List` type former.

Assume a type of shapes $S : \text{Type}$ and a type family of positions $P : S \rightarrow \text{Type}$. Assume that $S : \text{Type}$ is bridge-discrete, i.e. has an equivalence $\eta^S : (s_0 \equiv s_1) \simeq \text{Bridge}_S s_0 s_1$. Assume that $P : S \rightarrow \text{Type}$ is dependently bridge-discrete, that is, for every $s_0, s_1, (sr : s_0 \equiv s_1), (ss : \text{Bridge}_S s_0 s_1)$ such that $sr[\eta^S]ss$ and for every $p_0 : P s_0$ and $p_1 : P s_1$, there exists an equivalence $\eta^P : \text{PathP}_{i, P(sr\ i)} p_0 p_1 \simeq \text{BridgeP}_{X, P(ss\ x)} p_0 p_1$. Define $F : \text{Type} \rightarrow \text{Type}$ as $F X = \Sigma[s \in S] P s \rightarrow X$. Additionally define the following Agda data type μF :

```
data μF : Type where
  fold : F ( μF ) → μF
```

Note here that the data type declaration is accepted by Agda since F is a container functor and its input X occurs strictly positively⁸. We assert that the following equivalence holds $\mu F \simeq (X : \text{Type}) \rightarrow (F X \rightarrow X) \rightarrow X$. The proof follows a standard pattern. Recall that μF has an elimination principle $\mu Frec : \forall T \rightarrow (F T \rightarrow T) \rightarrow \mu F \rightarrow T$. Going from left to right we can define a map `toCh` using the latter principle. Going from right to left we can define a map $\lambda p \rightarrow p \mu F \text{fold}$. Proving the first inverse condition is done by induction. The other condition requires parametricity and reads as follows (using `funExt` whenever needed):

$$(p : (X : \text{Type}) \rightarrow (F X \rightarrow X) \rightarrow X) (A : \text{Type}) (f : F A \rightarrow A) \rightarrow \\ \text{toCh}(p \mu F \text{fold}) A f \equiv_A p A f$$

This of course looks like a global free theorem in the sense of (4). We wish to obtain this equality by applying `param` at program p and at the (logical) relation given by the graph of the function $\mu Frec A f : \mu F \rightarrow A$. We denote this graph as $\text{Gr}(\mu Frec A f) : \mu F \rightarrow A \rightarrow \text{Type}$. To that end we must supply a RRG structure for the domain of p and a dRRG structure for its codomain. For its domain we use `Tyfm` as above. For its codomain we must show $X : \text{Type} \models (F X \rightarrow X) \rightarrow X$ dRRG. Applying the `IIfm` rule of ROTT and other structural rules not displayed in Fig. 4 the goal is reduced to $X : \text{Type} \models F X$ dRRG. Recalling that $F X = \Sigma[s \in S] P s \rightarrow X$ we can apply the `Σfm` and `IIfm` rules of ROTT thereby reducing the goal to providing a RRG structure for S and providing $s : S \models P s$ dRRG. We can repackage our bridge-discreteness hypotheses η^S, η^P as such structures. At this stage all premises of `param` have been supplied and so we expect to obtain from `param` something of type $(\lambda X. (F X \rightarrow X) \rightarrow X) \{p \mu F, p A\}_{\text{Gr}(\mu Frec A f)}$. Again, contrary to its `BridgeP`

⁸Agda conveniently looks through the definition of F to decide this.

counterpart, the latter type computes to the appropriate relational parametricity statement, i.e.:

$$\text{param } \dots \mu F A (\text{Gr}(\mu \text{Frec } A f)) : \\ (f_0 : F \mu F \rightarrow \mu F)(f_1 : F A \rightarrow A)(fr : \dots) \rightarrow \mu \text{Frec } A f (p \mu F f_0) \equiv_A p A f_1$$

By setting $f_0 = \text{fold}$, $f_1 = f$ and providing an easy proof of their logical relatedness fr we get $\mu \text{Frec } A f (p \mu F \text{fold}) \equiv_A p A f$ and this proves the theorem.

Up to some reordering lemmas, we can obtain a Church encoding for the **List** data type as an instance of the above scheme of Church encodings: $\text{List } A \simeq (X : \text{Type}) \rightarrow X \rightarrow (A \rightarrow X \rightarrow X) \rightarrow X$. To do this the S and P parameters of the scheme are set to $S = 1 + A$ and $P(\text{inl } tt) = \perp$, $P(\text{inr } a) = \text{Unit}$. For this simpler **List** Church encoding, our bridge-discreteness hypotheses about S and P translate into the fact that A must be bridge-discrete for the encoding to hold. The reason is that, if A is not bridge-discrete, additional programs using their type variable non-parametrically will exist in the encoding. For instance **Type** is not bridge-discrete as its bridges are relations between types. Accordingly the encoding for **List** A with $A = \text{Type}$ does not hold, essentially because some polymorphic programs can use their type variable non parametrically: $\lambda X \text{ nl } cs X \text{ nl} : (X : \text{Type}) \rightarrow X \rightarrow (\text{Type} \rightarrow X \rightarrow X) \rightarrow X$. Similar considerations appear in [Nuyts and Devriese 2018].

4.3 System F

We can prove Reynolds's abstraction theorem [Reynolds 1983] for predicative System F [Leivant 1991] using ROTT and **param**. Indeed predicative System F is a subset of ROTT and the **param** theorem for dRRGs in that subset exactly expresses the abstraction theorem.

COROLLARY 4.1. *By analogy to Section 3.3.2, we have the following “inference rule”, which states that, given a kinding context Γ of predicative System F (consisting of type variables labeled with levels) and a type T of predicative System F over this context, all external dependent functions (i.e. all functions definable in Agda --bridges) from $\llbracket \Gamma \rrbracket$ to $\llbracket T \rrbracket$ respect logical relations, where $\llbracket - \rrbracket$ is an object-level translation from System F contexts (resp. types) to RRGs (resp. dRRGs).*

$$\frac{\Gamma \text{ Ctx-F} \quad \Gamma \vdash_F T \text{ type-F} \quad p : (\gamma : \llbracket \Gamma \rrbracket) \rightarrow \llbracket T \rrbracket_\gamma \quad \gamma_0, \gamma_1 : \llbracket \Gamma \rrbracket \quad \gamma r : \llbracket \Gamma \rrbracket \{ \gamma_0, \gamma_1 \}}{\text{param } \llbracket \Gamma \rrbracket \llbracket T \rrbracket p \gamma_0 \gamma_1 \gamma r : \llbracket T \rrbracket \{ p \gamma_0, p \gamma_1 \}_{\gamma r}} \text{PARAM-F}$$

We emphasize that this result proves parametricity of the *obvious* embedding $\llbracket - \rrbracket$ of predicative System F into Agda --bridges, which is definable in Agda --bridges. That is, $\llbracket - \rrbracket$ is defined the expected way. For instance, the System F type $X : *_0 \vdash_F X \rightarrow X$ type-F interprets as the dRRG $\text{Type}_0 \models X \rightarrow X$ dRRG which has $\lambda X. X \rightarrow X$ as its carrier. In other words, the dRRG model $\llbracket - \rrbracket$ is not some contrived construction, but simply a proof that all Agda --bridges types that read as System F types satisfy the SRP.

Hence we recover Reynolds's original notion of parametricity in Agda --bridges. We remark that [Nuyts et al. 2017] could not prove this result as it lacked the power of the **extent** rule and therefore could not properly characterize bridges between functions. Cavallo and Harper's [2021] system (which is almost identical to Agda --bridges) could prove this result equally well but the authors did not do this, and the same holds almost certainly for Bernardy et al. [2015]. Thus, as far as we are aware, this establishes the first formal relation between traditional parametricity as defined by Reynolds for System F using a logical relation, and internally parametric type theory.

5 REIMPLEMENTING HCOMP AND TRANSP

Recall from Section 2.1 that when compared to plain dependent type theory, cubical type theories feature additional language primitives: (1) a path interval, path types, path abstraction and application, (2) Kan operations, (3) additional type formers for turning equivalences into paths

in the universe, necessary to prove univalence, (4) (optionally) higher inductive types with path constructors.

The Kan operations of cubical type theory are necessary to make paths act as well-behaved proofs of equality. For instance, these operations are used to compose paths (i.e. prove transitivity of \equiv) or to turn the univalence map into an equivalence. The operational semantics of these operations is somewhat peculiar as it is defined by specifying how these operations reduce *at each type former*, i.e., by induction on the syntax of types. Hence the process of extending cubical type theory with a new type former F requires extra work: specifying how the Kan operations reduce at F .

Agda --bridges is an extension of Agda --cubical which implements CCHM cubical type theory [Cohen et al. 2017]. Thus, Agda --bridges must implement reduction clauses for the Kan operations of Agda --cubical at the **BridgeP** and **Gel** type formers. In Agda --cubical, the Kan operations are primitives named *homogeneous composition* (**hcomp**) and *transport* (**transp**).

```
transp : (A : I → Type) (φ : I) (u₀ : A i0) → A i1
hcomp : {A : Type} {φ : I} (u : (i : I) → Partial φ A) (u₀ : A) → A
```

We now explain the **transp** (Section 5.1) and **hcomp** (Section 5.2) operations and how Agda --bridges extend these. The exact equations can be consulted in our implementation. Further details about **transp**, **hcomp** can be found in [Vezzosi et al. 2021]. This section ends with a brief comparison between the Kan operations implemented by Agda --bridges and those specified by the CH theory.

5.1 Transport

The **transp** operation provides a way to coerce elements of a type into elements of a path-equal type. Indeed, given $AA : A_0 \equiv_{\text{Type}} A_1$ we obtain **transp** $(\lambda i. AA\ i)\ i0 : A_0 \rightarrow A_1$. This map can in turn be upgraded into an equivalence and thus **transp** can be used to validate one direction of the univalence equivalence $A_0 \equiv A_1 \rightarrow A_0 \simeq A_1$.

Regarding the $\varphi : I$ argument of **transp**, we merely indicate that it controls where the resulting coercion function is definitionally the identity. In other words, **transp** $A\ i1\ u_0$ reduces to u_0 (a typechecking side condition asks that A is constant when $\varphi = i1$). “Normal” transport is recovered by setting $\varphi = i0$ as illustrated above.

As explained before, the **transp** primitive reduces by induction on the formation of the line A , that is, a clause describing how **transp** reduces is specified when A is a line of Π -types, Σ -types, data types, record types, etc. Compared to --cubical, Agda --bridges implements two additional clauses for **transp**, handling lines of **BridgeP** and **Gel** types. For **Gel**, we indicate that the clause uses capturing (see Definition 2.2). We provide some details regarding the clause for **BridgeP** since it requires a generalized **hcomp** operation called **mhcomp** in Agda --bridges and explained hereafter.

Transporting bridges. Assume $A : I \rightarrow (@x : B I) \rightarrow \text{Type}$ and $a_\varepsilon : (i : I) \rightarrow A\ i\ \text{bie}$. The Agda --bridges implementation adds a clause for **transp** specifying how the following term should reduce **transp** $(i. \text{BridgeP}_{x.A\ i\ x}\ (a_0\ i)\ (a_1\ i))\ (\varphi : I)\ u_0$. This term can be represented as the dotted line in the following diagram.

$$\begin{array}{ccc}
 a_0\ i1 & \text{-----} & a_1\ i1 \\
 \uparrow & & \uparrow \\
 a_0\ i0 & \xrightarrow{u_0} & a_1\ i0
 \end{array}$$

Setting the result of this reduction to be the bridge $\lambda(x : B I). \text{transp}\ (i. A\ i\ x)\ \varphi\ (u_0\ x)$ does not work. Indeed the latter term does not have the required endpoints $a_0\ i1, a_1\ i1$ definitionally.

A second idea is to use the **hcomp** primitive (its type appears above) of --cubical, which is used to reduce **transp** along lines of path types. In general, the sole purpose of **hcomp** is to allow terms to be

definitionally adjusted on a fragment of the type they live in. More precisely, suppose the ambient context Γ contains a number of path variables $\Phi = j, k, \dots$ and suppose $\Gamma \vdash \chi : \mathbf{l}$. Intuitively, the term $\Gamma \vdash \text{hcomp} \{B : \text{Type}\} \{\chi : \mathbf{l}\} v (v_0 : B)$ is $v_0 : B$ but definitionally adjusted (using the v argument, not explained here) on the fragment of $\Gamma \vdash B$ where $\chi(j, k, \dots) = \mathbf{i1}$ (since v_0, B, χ live in context Γ they can depend on variables in Φ). In cubical type theory, constraints like $\chi(j, k, \dots) = \mathbf{i1}$ are called *face constraints* or alternatively *cofibrations* and can be regarded as subsets of a cube $\chi \subseteq \Phi$. The language used to express face constraints is called a *face or cofibration logic*. The cofibration logic used by Agda --cubical is De Morgan algebra and assertions in this logic are encoded as terms $\chi : \mathbf{l}$. If $\Gamma \vdash j, k, l : \mathbf{l}$, an example of face constraint is $\chi = ((\sim j) \vee k) \wedge l$.

In our case we would like, in a context extended with $(x : \mathbf{Bl})$, to definitionally adjust the term $\text{transp} (i. A \ i \ x) \varphi (u_0 \ x)$ of type $A \ \mathbf{i1} \ x$ when $x = \mathbf{bi0}$ and $x = \mathbf{bi1}$ (and in fact when $\varphi = \mathbf{i1}$). The problem of course is that x is a bridge variable, and that hcomp only allows constraints $\varphi : \mathbf{l}$ on path variables. The *mixed homogeneous composition* mhcomp primitive of Agda --bridges generalizes hcomp w.r.t. bridge variables and can be used in place of hcomp to specify the result of transporting along a line of bridges.

5.2 Mixed Homogeneous Composition

The mhcomp primitive of Agda --bridges has a type different than that of hcomp .

$$\text{mhcomp} : \forall \{A : \text{Type}\} \{\zeta : \text{MCstr}\} (u : (i : \mathbf{l}) \rightarrow \text{MPartial} \zeta A) (u_0 : A) \rightarrow A$$

This time u_0 and its type A can have free path variables $\Phi = (j, k, \dots)$ but also free bridge variables $\Psi = (x, y, \dots)$ and u_0 ought to be definitionally adjusted on a subset ζ of the mixed cube $\Phi \times \Psi$. For this reason mhcomp expects face constraints ζ expressed in an extended cofibration logic called MCstr . Concretely, the latter is a type postulated by Agda --bridges and equipped with primitives for combining atomic mixed face constraints. Instead of precisely explaining these primitives we provide a formula $\text{MCstr}(\Phi, \Psi)$ expressing what mixed constraints ζ can be built in a context containing Φ and Ψ as above. First, define $\mathbf{l}(\Phi) = \{\varphi \mid \Phi \vdash \varphi : \mathbf{l}\}$. Second, define the set of bridge hyperfaces of Ψ as $H(\Psi) = \Psi \times \{\mathbf{bi0}, \mathbf{bi1}\}$. We define $\text{BCstr}(\Psi) = \{\bigvee_{(x, \mathbf{bie}) \in H'} (x = \mathbf{bie}) \mid H' \subseteq H\} \cup \{\top\}$, i.e., bridge face constraints obtainable in Ψ are disjunctions of bridge hyperfaces (this includes an empty disjunction \perp), or a vacuous constraint denoted \top . Finally we set

$$\text{MCstr}(\Phi, \Psi) = \frac{\mathbf{l}(\Phi) \times \text{BCstr}(\Psi)}{\forall \varphi \psi. (\mathbf{i1}, \psi) = (\varphi, \top) \Rightarrow \top_{\text{MCstr}}}$$

The quotient is taken to turn the map $\varphi \mapsto (\varphi, \perp)$ into an *embedding* of logics: a --cubical constraint $\varphi : \mathbf{l}$ holds if and only if its image $(\varphi, \perp) : \text{MCstr}$ is a mixed constraint that holds. This condition is required to ensure that a term typechecks in Agda --cubical if and only if it typechecks in Agda --bridges. An example of mixed constrained $\zeta : \text{MCstr}$ is $\zeta := (\varphi, (x = \mathbf{bi0}) \vee (x = \mathbf{bi1}))$ which appears when transporting bridges, as hinted above.

$$\begin{aligned} & \text{transp} (i. \text{BridgeP}_{x. A \ i \ x} (a_0 \ i) (a_1 \ i)) \varphi u_0 \mapsto \\ & \lambda(x : \mathbf{Bl}). \text{mhcomp} \{A \ \mathbf{i1} \ x\} \{(\varphi, (x = \mathbf{bi0}) \vee (x = \mathbf{bi1}))\} (\dots) (\text{transp} (i. A \ i \ x) \varphi (u_0 \ x)) \end{aligned}$$

Similar to transp , the operational semantics of mhcomp is defined by induction on the syntax of its $A : \text{Type}$ argument. Concretely, Agda --bridges duplicates the hcomp equations for Glue , hcomp , PathP , Σ , Π , record and (non-indexed, non-HIT) data types, but propagating a mixed constraint ζ this time. Additionally, it implements reductions at BridgeP and Gel types. The latter clause uses capturing. Following --cubical, if A is a HIT, an inhabitant $\text{mhcomp} \{A\} \{\zeta\} u u_0 : A$ is considered normal and functions defined by pattern matching on A compute on it if $\zeta = (\varphi, \perp)$ for some φ .

Comparison with the CH Theory. While the theory of Agda --bridges extends CCHM cubical type theory, the CH theory is an extension of cartesian cubical type theory [Angiuli et al. 2021a, 2018] (CCTT). This distinction leads to Kan operations formulated differently in each system. First, in order to validate univalence, CCTT postulates V-types whereas CCHM postulates *Glue* types. Hence *transp* and *mhcomp* both have a reduction clause for *Glue* types instead of V-types. Second, the composition Kan operation of CCTT uses a simple cofibration logic with a constructor for diagonal constraints ($x = y : I$). By contrast, CCHM uses a De Morgan cofibration logic which Agda --bridges had to extend (see *MCstr* above). To pinpoint a sound definition for *MCstr* we inspected the presheaf model $\text{psh}(\square_{\text{DM}} \times \square_a)$ where \square_{DM} (resp. \square_a) is the category of De Morgan cubes (see CCHM; resp. affine cubes, see CH). The above formula for *MCstr*(Φ, Ψ) can be regarded as the definition of such a presheaf. Finally some equations for the Kan operations use capturing, handled differently in Agda --bridges and CH: sound capturing is performed through context restriction in the CH theory (see Section 2.2) and through semi-freshness in Agda --bridges (see Section 2.4).

6 RELATED WORK

6.1 Relational Parametricity in and of Type Theory

In order to discuss and classify related work about parametricity in and of type theory, we analyze the general statement of parametricity. We assume that we are studying an object language \mathcal{X} which embeds or can be interpreted in a target language \mathcal{Y} where free theorems [Wadler 1989] will be stated.

Parametricity: For every (\mathcal{S}) type T in \mathcal{X} , there exists a logical relation $[T]$ in \mathcal{Y} , such that every (\mathcal{P}) program $t : T$ in \mathcal{X} is self-related according to $[T]$ in \mathcal{Y} .

Note that the general statement of parametricity contains two universal quantifications, which we have labeled with names \mathcal{S} and \mathcal{P} for the (meta)theories where these quantifications take place. We will classify related work on parametricity by its choice of languages/theories \mathcal{X} , \mathcal{Y} , \mathcal{P} and \mathcal{S} . Note that if $\mathcal{X} = \mathcal{Y} = \mathcal{P}$, then we have global free theorems in \mathcal{Y} and can prove in \mathcal{Y} e.g. that Church encodings in \mathcal{X} behave as the data type they encode. By contrast if \mathcal{P} is a metatheory, then we only have free theorems for concretely known programs and the framework is merely a parametricity *translation* of such programs. These situations are often referred to as *internal* vs. *external* parametricity. In any provenly sound framework for internal parametricity hitherto developed that we are aware of, \mathcal{S} is a metatheory (at least if you want $[T]$ to be a concrete relation and not a bridge type whose meaning needs to be characterized separately), implying that the SRP (Section 3.1.2) is a metatheorem.

Languages with internal parametricity typically feature non-standard parametricity primitives which call for a denotational model to establish soundness. In this case, when relevant, we specify how (closed) types in \mathcal{Y} are modelled, and the metatheory \mathcal{M} in which the model is built. In the case of external parametricity, we rather specify how \mathcal{X} is interpreted in the metatheory \mathcal{Y} .

The resulting classification is given in Table 1 (Ab denotes Agda --bridges). The first block lists treatments of System F. Reynolds's original model is in set theory, but was later shown not to support impredicativity [Reynolds 1984]. Subsequent treatments use a logic over System F and prove parametricity results about concrete programs of concrete types. Atkey explains how Reynolds's model can be restructured as a reflexive graph model and then generalizes to System F ω . The third block lists treatments of external parametricity for dependent types. The first three papers prove parametricity results again in dependent type systems, but only for concrete programs of concrete types. The latter two use denotational models. Tabareau et al.'s approach is peculiar in that it defines the logical relation on the universe as the type of relations *that are the graph of an equivalence*, thus establishing a framework not for relational but for univalent parametricity.

Table 1. Classification of related work about parametricity, based on [Nuyts et al. 2017, fig. 9]. Here, \mathcal{Y} is (faintly) highlighted when (possibly) $\mathcal{Y} = \mathcal{X}$; \mathcal{P} is highlighted if $\mathcal{P} = \mathcal{Y} = \mathcal{X}$; and \mathcal{S} is highlighted if $\mathcal{S} = \mathcal{X}$.

Citation	Obj. lang. \mathcal{X}	Target lang. \mathcal{Y}	\mathcal{P}	\mathcal{S}	\mathcal{M}	model in \mathcal{Y} or (<i>italic</i>) \mathcal{M}
[Reynolds 1983] [Abadi et al. 1993] [Plotkin and Abadi 1993] [Wadler 2007]	System F System F System F System F	Meta: Set theory System \mathcal{R} System F + logic System F + logic	\mathcal{Y} Meta Meta Meta	\mathcal{Y} Meta Meta Meta	Meta	Sets with relations PERs [Bellucci et al. 1995]
[Atkey 2012]	System F ω	Meta: Impred. CIC	\mathcal{Y}	\mathcal{Y}		Reflexive graphs
[Takeuti 2001] [Bernardy et al. 2012] [Tabareau et al. 2021] [Krishnaswami and Dreyer 2013] [Atkey et al. 2014]	$\mathcal{X} \in \lambda$ -cube Any PTS CIC CC MLTT	$\mathcal{X} + \forall + \Pi \in \lambda$ -cube Suitable PTS Univalent CIC Meta Meta: CIC	Meta Meta Meta Meta \mathcal{Y}	Meta Meta Meta Meta \mathcal{Y}		PERs Reflexive graphs
[Bernardy and Moulin 2012] [Bernardy et al. 2015] [Nuyts et al. 2017] [Nuyts and Devriese 2018] [Cavallo and Harper 2021] Agda --bridges (Ab)	BM BCM ParamDTT RelDTT CH Ab	\mathcal{X} \mathcal{X} \mathcal{X} \mathcal{X} \mathcal{X} \mathcal{X}	\mathcal{X} \mathcal{X} \mathcal{X} \mathcal{X} \mathcal{X} \mathcal{X}	\mathcal{X} Meta Meta Meta Meta Meta	Meta Meta Meta Meta Meta Meta	<i>none</i> Affine cubical sets Bridge/path cubical sets Depth n cubical sets Affine/cart. bicub. sets Affine/CCHM bicub. sets
ROTT Corollary 4.1	DTT (\sim Ab) Pr. Sys. F	Ab Ab	\mathcal{Y} \mathcal{Y}	\mathcal{Y} \mathcal{Y}		RRGs RRGs
[Altenkirch et al. 2024] Cub. ROTT (envisioned)	ACKS \approx ACKS	\mathcal{X} \mathcal{X}	\mathcal{X} \mathcal{X}	\mathcal{X} \mathcal{X}	Meta Ab	Affine cubical sets Relativistic cubical sets

The fourth block lists treatments of internal parametricity for dependent types; these produce concrete free theorems (i.e. not mentioning bridge types that need to be characterized separately) for abstract programs of concrete types. Bernardy and Moulin’s [2012] system predates the usage of named relational dimensions, but it is observational, i.e. the SRP holds definitionally. We are unaware of any soundness proof for this system. Bernardy et al. introduce the bridge interval as well as the *extent* and *Gel* combinators, and Cavallo and Harper combine their system with HoTT and demonstrate its power on paper. With Agda --bridges, we provide an implementation. The work on ParamDTT and RelDTT introduces a modal system in order to prove Reynolds’s identity extension lemma for large types, but in the process has to adopt a cartesian cubical model which does not validate *extent*. As a consequence, these systems lack the power to prove parametricity of System F, which we demonstrated can be done in Agda --bridges (Section 4.3). We refer to Nuyts [2021] for a brief discussion of various internal parametricity features and their requirements in the model.

Strictly speaking, ROTT is again a dependently typed system with external parametricity that could go in the third block. However, ROTT was conceived as a commodity for Agda --bridges and seeks to obtain free theorems there. This is in contrast with e.g. Atkey et al. [2014], where MLTT is the system of interest but free theorems are obtained in some metatheory. We remark that since *param* is an external rule, the source syntax of ROTT is really just dependent type theory (extensible with bridge types, which are also displayed RRGs).

The reason why ROTT cannot provide internal parametricity is that the logical relation specified for an RRG, is itself not a displayed RRG (i.e. a dependent ROTT type) but only an external Agda --bridges type. This might be remedied in the future by moving from RRGs to relativistic cubical types, i.e. cubical types whose n -cubes are equivalent to n -cubes of bridges. Such a system would allow internal parametricity and satisfy the SRP definitionally. The syntax of such a system would

be very close to that of the concurrent work by [Altenkirch et al. \[2024\]](#), who present an internally parametric type theory which avoids the use of an interval and validates the SRP definitionally.

6.2 Parametricity and Univalence

While relational parametricity requires functions to respect relations, HoTT asks that they respect equivalences. Of course, equivalences are a form of relations, so one idea is akin to the other.

We already mentioned [Tabareau et al.’s \[2021\]](#) work in the previous section. A reformulation of their work using our techniques would effectively lead to a shallow embedding of observational setoid [\[Pujet and Tabareau 2022\]](#) or homotopy [\[Altenkirch et al. 2022\]](#) type theory in a CwF of univalent setoids or groupoids.

[Awodey et al. \[2018\]](#) define Church-like encodings of ordinary but also higher inductive types in (homotopy) type theory. Since the correctness of the Church encoding relies on preservation not only of equivalences but of all relations, they cannot be proven correct in plain type theory or HoTT. Instead, the authors enforce relational parametricity simply by adding it as a “such that” clause to the encoding.

6.3 The Structure Identity Principle (SIP)

We discuss appearances of the SIP in the literature and compare these to our discussion in Section 3. The HoTT book does not actually feature the full SIP and two other treatments rely on a DSL.

Standard notions of structure on univalent categories. The HoTT book [\[Program 2013\]](#) defines a *notion of structure* on a category C essentially as a displayed category \mathcal{D}^\bullet over C such that the projection functor $P : \Sigma C \mathcal{D}^\bullet \rightarrow C$ from the total space is faithful, implying that the fiber of any object of C (and its identity morphism) is a preorder. A notion of structure is *standard* if this preorder is always a partial order, which is defined as a univalent preorder. The theorem called SIP then states that if C is univalent and \mathcal{D}^\bullet is standard, then the total space $\Sigma C \mathcal{D}^\bullet$ is univalent. Relevant examples are group structures over h-sets, setoid structures over h-sets, monad structures over endofunctors, functor structures over indexed objects, ...

Fundamentally, the proof of this theorem does two things: It applies the extensionality principle for Σ -types to characterize a path between objects in $\Sigma C \mathcal{D}^\bullet$, and it uses path induction to deduce “displayed” univalence from standardness (which meant fiberwise univalence). Importantly, if \mathcal{D}^\bullet is quite complex, it is still up to the user to prove fiberwise univalence there, which still requires either rote work or the usage of a DSL for standard notions of structure. As such, we see this SIP as only a fragment of the fully general SIP.

A DSL for univalent structures on Type. [Angiuli et al. \[2021b\]](#) are concerned with proving the SIP (in our most general sense) for types of the form $T = \Sigma[X : \text{Type}] \Sigma[s : SX] P X s$, where $P X s$ is a mere proposition (h-prop). The idea is that a tuple (X, s, p) is an algebra-like object with carrier X and operations s satisfying the axioms $P X s$. Their paper features a theorem titled SIP which amounts to the characterization of paths in a type of the form $\Sigma[X : \text{Type}] C X$. Since paths in a mere proposition can be characterized as informationless (Theorem 3.6), only the SIP for the operations type SX remains to be dealt with. A DSL – essentially the type language of the STLC with base type X (the carrier) – is then provided to build structures satisfying the SIP.

A DSL for univalent higher categories. [Ahrens et al. \[2020, thm. 7.10\]](#) use FOLDS [\[Makkai 1995\]](#) as a DSL for building univalent higher categories. In other words, they show that a higher type satisfies the SIP if it occurs as the object type of a higher category specified by a FOLDS signature.

ACKNOWLEDGMENTS

We thank Andrea Vezzosi for continuously sharing with us his expertise and sound suggestions regarding Agda --cubical and Agda --bridges. We thank Rasmus Møgelberg and Andrea Vezzosi for welcoming the first author at the IT University of Copenhagen. We thank the reviewers for their remarks and suggestions. Antoine Van Muylder holds a PhD fellowship (11H9921N) of the Research Foundation – Flanders (FWO). Andreas Nuyts holds a Postdoctoral fellowship (1247922N) of the Research Foundation – Flanders (FWO). This research is partially funded by the Research Fund KU Leuven and by the Research Foundation - Flanders (FWO; G030320N).

DATA AVAILABILITY STATEMENT

We provide in [Van Muylder et al. 2023] a virtual machine where Agda --bridges is installed and where our Agda --bridges library and its results are typechecked. The Agda --bridges implementation is developed in <https://github.com/antoinevanmuylder/agda/tree/bridges> and our library can be found in <https://github.com/antoinevanmuylder/bridgy-lib>. For some documentation, see the latter repository. Note that some notions have slightly different names than in this work.

REFERENCES

- Martín Abadi, Luca Cardelli, and Pierre-Louis Curien. 1993. Formal Parametric Polymorphism. *Theor. Comput. Sci.* 121, 1&2 (1993), 9–58. [https://doi.org/10.1016/0304-3975\(93\)90082-5](https://doi.org/10.1016/0304-3975(93)90082-5)
- Michael Gordon Abbott, Thorsten Altenkirch, and Neil Ghani. 2005. Containers: Constructing strictly positive types. *Theor. Comput. Sci.* 342, 1 (2005), 3–27. <https://doi.org/10.1016/J.TCS.2005.06.002>
- The Agda Community. [n. d.]. A standard library for Cubical Agda. <https://github.com/agda/cubical>
- Agda Development Team. 2023. Agda 2.6.3 documentation. <https://agda.readthedocs.io/en/v2.6.3/>
- Benedikt Ahrens and Peter LeFanu Lumsdaine. 2019. Displayed Categories. *Log. Methods Comput. Sci.* 15, 1 (2019). [https://doi.org/10.23638/LMCS-15\(1:20\)2019](https://doi.org/10.23638/LMCS-15(1:20)2019)
- Benedikt Ahrens, Paige Randall North, Michael Shulman, and Dimitris Tsementzis. 2020. A Higher Structure Identity Principle. In *LICS '20: 35th Annual ACM/IEEE Symposium on Logic in Computer Science, Saarbrücken, Germany, July 8-11, 2020*, Holger Hermanns, Lijun Zhang, Naoki Kobayashi, and Dale Miller (Eds.). ACM, 53–66. <https://doi.org/10.1145/3373718.3394755>
- Thorsten Altenkirch, Yorgo Chamoun, Ambrus Kaposi, and Michael Shulman. 2024. Internal parametricity, without an interval. In *To appear in: Proceedings of the 51st Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2024*. ACM.
- Thorsten Altenkirch and Ambrus Kaposi. 2015. Towards a Cubical Type Theory without an Interval. In *21st International Conference on Types for Proofs and Programs, TYPES 2015, May 18-21, 2015, Tallinn, Estonia (LIPIcs, Vol. 69)*, Tarmo Uustalu (Ed.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 3:1–3:27. <https://doi.org/10.4230/LIPICS.TYPES.2015.3>
- Thorsten Altenkirch, Ambrus Kaposi, and Michael Shulman. 2022. Towards Higher Observational Type Theory. In *28th International Conference on Types for Proofs and Programs (TYPES 2022)*. University of Nantes.
- Thorsten Altenkirch, Conor McBride, and Wouter Swierstra. 2007. Observational equality, now!. In *Proceedings of the ACM Workshop Programming Languages meets Program Verification, PLPV 2007, Freiburg, Germany, October 5, 2007*, Aaron Stump and Hongwei Xi (Eds.). ACM, 57–68. <https://doi.org/10.1145/1292597.1292608>
- Abhishek Anand and Greg Morrisett. 2017. Revisiting Parametricity: Inductives and Uniformity of Propositions. *CoRR* abs/1705.01163 (2017). arXiv:1705.01163 <http://arxiv.org/abs/1705.01163>
- Carlo Angiuli, Guillaume Brunerie, Thierry Coquand, Robert Harper, Kuen-Bang Hou (Favonia), and Daniel R. Licata. 2021a. Syntax and models of Cartesian cubical type theory. *Math. Struct. Comput. Sci.* 31, 4 (2021), 424–468. <https://doi.org/10.1017/S0960129521000347>
- Carlo Angiuli, Evan Cavallo, Anders Mörtberg, and Max Zeuner. 2021b. Internalizing representation independence with univalence. *Proc. ACM Program. Lang.* 5, POPL (2021), 1–30. <https://doi.org/10.1145/3434293>
- Carlo Angiuli, Kuen-Bang Hou (Favonia), and Robert Harper. 2018. Cartesian Cubical Computational Type Theory: Constructive Reasoning with Paths and Equalities. In *27th EACSL Annual Conference on Computer Science Logic, CSL 2018, September 4-7, 2018, Birmingham, UK (LIPIcs, Vol. 119)*, Dan R. Ghica and Achim Jung (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 6:1–6:17. <https://doi.org/10.4230/LIPICS.CSL.2018.6>
- Robert Atkey. 2012. Relational Parametricity for Higher Kinds. In *Computer Science Logic (CSL '12) - 26th International Workshop/21st Annual Conference of the EACSL, CSL 2012, September 3-6, 2012, Fontainebleau, France (LIPIcs, Vol. 16)*, Patrick Cégielski and Arnaud Durand (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 46–61. <https://doi.org/>

10.4230/LIPICS.CSL.2012.46

- Robert Atkey, Neil Ghani, and Patricia Johann. 2014. A relationally parametric model of dependent type theory. In *The 41st Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '14, San Diego, CA, USA, January 20-21, 2014*, Suresh Jagannathan and Peter Sewell (Eds.). ACM, 503–516. <https://doi.org/10.1145/2535838.2535852>
- Steve Awodey, Jonas Frey, and Sam Speight. 2018. Impredicative Encodings of (Higher) Inductive Types. In *Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science, LICS 2018, Oxford, UK, July 09-12, 2018*, Anuj Dawar and Erich Grädel (Eds.). ACM, 76–85. <https://doi.org/10.1145/3209108.3209130>
- Roberto Bellucci, Martín Abadi, and Pierre-Louis Curien. 1995. A Model for Formal Parametric Polymorphism: A PER Interpretation for System R. In *Typed Lambda Calculi and Applications, Second International Conference on Typed Lambda Calculi and Applications, TLCA '95, Edinburgh, UK, April 10-12, 1995, Proceedings (Lecture Notes in Computer Science, Vol. 902)*, Mariangiola Dezani-Ciancaglini and Gordon D. Plotkin (Eds.). Springer, 32–46. <https://doi.org/10.1007/BFB0014043>
- Jean-Philippe Bernardy, Thierry Coquand, and Guilhem Moulin. 2015. A Presheaf Model of Parametric Type Theory. In *The 31st Conference on the Mathematical Foundations of Programming Semantics, MFPS 2015, Nijmegen, The Netherlands, June 22-25, 2015 (Electronic Notes in Theoretical Computer Science, Vol. 319)*, Dan R. Ghica (Ed.). Elsevier, 67–82. <https://doi.org/10.1016/J.ENTCS.2015.12.006>
- Jean-Philippe Bernardy, Patrik Jansson, and Ross Paterson. 2012. Proofs for free - Parametricity for dependent types. *J. Funct. Program.* 22, 2 (2012), 107–152. <https://doi.org/10.1017/S0956796812000056>
- Jean-Philippe Bernardy and Guilhem Moulin. 2012. A Computational Interpretation of Parametricity. In *Proceedings of the 27th Annual IEEE Symposium on Logic in Computer Science, LICS 2012, Dubrovnik, Croatia, June 25-28, 2012*. IEEE Computer Society, 135–144. <https://doi.org/10.1109/LICS.2012.25>
- Auke Bart Booi, Martín Hötzel Escardó, Peter LeFanu Lumsdaine, and Michael Shulman. 2016. Parametricity, Automorphisms of the Universe, and Excluded Middle. In *22nd International Conference on Types for Proofs and Programs, TYPES 2016, May 23-26, 2016, Novi Sad, Serbia (LIPIcs, Vol. 97)*, Silvia Ghilezan, Herman Geuvers, and Jelena Ivetic (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 7:1–7:14. <https://doi.org/10.4230/LIPICS.TYPES.2016.7>
- Simon Boulrier, Pierre-Marie Pédro, and Nicolas Tabareau. 2017. The next 700 syntactical models of type theory. In *Proceedings of the 6th ACM SIGPLAN Conference on Certified Programs and Proofs, CPP 2017, Paris, France, January 16-17, 2017*, Yves Bertot and Viktor Vafeiadis (Eds.). ACM, 182–194. <https://doi.org/10.1145/3018610.3018620>
- Evan Cavallo. 2020. *Ptt, an experimental implementation of Martin-Löf type theory with n-ary internal parametricity*. <https://github.com/ecavallo/ptt>
- Evan Cavallo and Robert Harper. 2019. Parametric Cubical Type Theory. *CoRR* abs/1901.00489 (2019). arXiv:1901.00489 <http://arxiv.org/abs/1901.00489>
- Evan Cavallo and Robert Harper. 2021. Internal Parametricity for Cubical Type Theory. *Log. Methods Comput. Sci.* 17, 4 (2021). [https://doi.org/10.46298/LMCS-17\(4:5\)2021](https://doi.org/10.46298/LMCS-17(4:5)2021)
- Cyril Cohen, Thierry Coquand, Simon Huber, and Anders Mörtberg. 2017. Cubical Type Theory: A Constructive Interpretation of the Univalence Axiom. *FLAP* 4, 10 (2017), 3127–3170. <http://collegepublications.co.uk/ifcolog/200019>
- Jean-Yves Girard. 1986. The System F of Variable Types, Fifteen Years Later. *Theor. Comput. Sci.* 45, 2 (1986), 159–192. [https://doi.org/10.1016/0304-3975\(86\)90044-7](https://doi.org/10.1016/0304-3975(86)90044-7)
- Jean-Yves Girard. 1972. *Interprétation fonctionnelle et élimination des coupures de l'arithmétique d'ordre supérieur*. Ph.D. Dissertation. Éditeur inconnu.
- Claudio Hermida, Uday S. Reddy, and Edmund P. Robinson. 2013. Logical Relations and Parametricity - A Reynolds Programme for Category Theory and Programming Languages. In *Proceedings of the Workshop on Algebra, Coalgebra and Topology, WACT 2013, Bath, UK, March 1, 2013 (Electronic Notes in Theoretical Computer Science, Vol. 303)*, John Power and Cai Wingfield (Eds.). Elsevier, 149–180. <https://doi.org/10.1016/J.ENTCS.2014.02.008>
- Chantal Keller and Marc Lasson. 2012. Parametricity in an Impredicative Sort. In *Computer Science Logic (CSL'12) - 26th International Workshop/21st Annual Conference of the EACSL, CSL 2012, September 3-6, 2012, Fontainebleau, France (LIPIcs, Vol. 16)*, Patrick Cégielski and Arnaud Durand (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 381–395. <https://doi.org/10.4230/LIPICS.CSL.2012.381>
- Neelakantan R. Krishnaswami and Derek Dreyer. 2013. Internalizing Relational Parametricity in the Extensional Calculus of Constructions. In *Computer Science Logic 2013 (CSL 2013), CSL 2013, September 2-5, 2013, Torino, Italy (LIPIcs, Vol. 23)*, Simona Ronchi Della Rocca (Ed.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 432–451. <https://doi.org/10.4230/LIPICS.CSL.2013.432>
- Daniel Leivant. 1991. Finitely Stratified Polymorphism. *Inf. Comput.* 93, 1 (1991), 93–113. [https://doi.org/10.1016/0890-5401\(91\)90053-5](https://doi.org/10.1016/0890-5401(91)90053-5)
- Michael Makkai. 1995. First order logic with dependent sorts, with applications to category theory. (1995). <http://www.math.mcgill.ca/makkai/folds/foldsinpdf/FOLDS.pdf>

- Bassel Manna and Rasmus Ejlers Møgelberg. 2018. The Clocks They Are Adjunctions Denotational Semantics for Clocked Type Theory. In *3rd International Conference on Formal Structures for Computation and Deduction, FSCD 2018, July 9-12, 2018, Oxford, UK (LIPIcs, Vol. 108)*, Hélène Kirchner (Ed.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 23:1–23:17. <https://doi.org/10.4230/LIPIcs.FSCD.2018.23>
- Guilhem Moulin. 2016. *Internalizing Parametricity*. Ph. D. Dissertation. Chalmers University of Technology, Gothenburg, Sweden. <http://publications.lib.chalmers.se/publication/235758-internalizing-parametricity>
- Andreas Nuyts. 2021. Parametricity Features and their Requirements. *CoRR* abs/2111.09822 (2021). arXiv:2111.09822 <https://arxiv.org/abs/2111.09822>
- Andreas Nuyts and Dominique Devriese. 2018. Degrees of Relatedness: A Unified Framework for Parametricity, Irrelevance, Ad Hoc Polymorphism, Intersections, Unions and Algebra in Dependent Type Theory. In *Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science, LICS 2018, Oxford, UK, July 09-12, 2018*, Anuj Dawar and Erich Grädel (Eds.). ACM, 779–788. <https://doi.org/10.1145/3209108.3209119>
- Andreas Nuyts, Andrea Vezzosi, and Dominique Devriese. 2017. Parametric quantifiers for dependent type theory. *Proc. ACM Program. Lang.* 1, ICFP (2017), 32:1–32:29. <https://doi.org/10.1145/3110276>
- Gordon D. Plotkin and Martín Abadi. 1993. A Logic for Parametric Polymorphism. In *Typed Lambda Calculi and Applications, International Conference on Typed Lambda Calculi and Applications, TLCA '93, Utrecht, The Netherlands, March 16-18, 1993, Proceedings (Lecture Notes in Computer Science, Vol. 664)*, Marc Bezem and Jan Friso Groote (Eds.). Springer, 361–375. <https://doi.org/10.1007/BFB0037118>
- The Univalent Foundations Program. 2013. *Homotopy Type Theory: Univalent Foundations of Mathematics*. Institute for Advanced Study. <https://homotopytypetheory.org/book/>
- Loïc Pujet and Nicolas Tabareau. 2022. Observational equality: now for good. *Proc. ACM Program. Lang.* 6, POPL (2022), 1–27. <https://doi.org/10.1145/3498693>
- John C. Reynolds. 1974. Towards a theory of type structure. In *Programming Symposium, Proceedings Colloque sur la Programmation, Paris, France, April 9-11, 1974 (Lecture Notes in Computer Science, Vol. 19)*, Bernard J. Robinet (Ed.). Springer, 408–423. https://doi.org/10.1007/3-540-06859-7_148
- John C. Reynolds. 1983. Types, Abstraction and Parametric Polymorphism. In *Information Processing 83, Proceedings of the IFIP 9th World Computer Congress, Paris, France, September 19-23, 1983*, R. E. A. Mason (Ed.). North-Holland/IFIP, 513–523.
- John C. Reynolds. 1984. Polymorphism is not Set-Theoretic. In *Semantics of Data Types, International Symposium, Sophia-Antipolis, France, June 27-29, 1984, Proceedings (Lecture Notes in Computer Science, Vol. 173)*, Gilles Kahn, David B. MacQueen, and Gordon D. Plotkin (Eds.). Springer, 145–156. https://doi.org/10.1007/3-540-13346-1_7
- Egbert Rijke. 2022. Introduction to Homotopy Type Theory. arXiv:2212.11082 [math.LO]
- Robert Rose, Matthew Z Weaver, and Daniel R Licata. 2022. Deciding the cofibration logic of cartesian cubical type theories. In *28th International Conference on Types for Proofs and Programs (TYPES 2022)*. University of Nantes.
- Nicolas Tabareau, Éric Tanter, and Matthieu Sozeau. 2021. The Marriage of Univalence and Parametricity. *J. ACM* 68, 1, Article 5 (jan 2021), 44 pages. <https://doi.org/10.1145/3429979>
- Izumi Takeuti. 2001. The Theory of Parametricity in Lambda Cube. Technical report 1217, Kyoto University. https://repository.kulib.kyoto-u.ac.jp/dspace/bitstream/2433/41237/1/1217_10.pdf
- The Coq development team. 2022. The Coq proof assistant. <http://coq.inria.fr>
- Antoine Van Muylder, Andreas Nuyts, and Dominique Devriese. 2023. *Agda --bridges VM*. <https://doi.org/10.5281/zenodo.10009365>
- Niccolò Veltri and Andrea Vezzosi. 2020. Formalizing π -calculus in guarded cubical Agda. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2020, New Orleans, LA, USA, January 20-21, 2020*, Jasmin Blanchette and Catalin Hritcu (Eds.). ACM, 270–283. <https://doi.org/10.1145/3372885.3373814>
- Niccolò Veltri and Andrea Vezzosi. 2023. Formalizing CCS and π -calculus in Guarded Cubical Agda. *J. Log. Algebraic Methods Program.* 131 (2023), 100846. <https://doi.org/10.1016/J.JLAMP.2022.100846>
- Andrea Vezzosi, Anders Mörtberg, and Andreas Abel. 2021. Cubical Agda: A dependently typed programming language with univalence and higher inductive types. *J. Funct. Program.* 31 (2021), e8. <https://doi.org/10.1017/S0956796821000034>
- Philip Wadler. 1989. Theorems for Free!. In *Proceedings of the fourth international conference on Functional programming languages and computer architecture, FPCA 1989, London, UK, September 11-13, 1989*, Joseph E. Stoy (Ed.). ACM, 347–359. <https://doi.org/10.1145/99370.99404>
- Philip Wadler. 2007. The Girard-Reynolds isomorphism (second edition). *Theor. Comput. Sci.* 375, 1-3 (2007), 201–226. <https://doi.org/10.1016/J.TCS.2006.12.042>

Received 2023-07-11; accepted 2023-11-07